The Worst Prediction in Physics Is Not a Prediction: Axiom Calibration of the Cosmological Constant Problem

A Lean 4 Formalization (Paper 42)

Paul Chun-Kit Lee*
New York University
dr.paul.c.lee@gmail.com

February 2026

Paper 42 of the Constructive Reverse Mathematics Programme

Abstract

We apply the axiom calibration framework of constructive reverse mathematics to the cosmological constant problem—the alleged 10^{120} discrepancy between quantum field theory's "prediction" of vacuum energy and the observed value. The problem decomposes into three logically distinct claims with different constructive status. (I) The 10^{120} ultraviolet discrepancy is a regulator-dependent artifact: it arises only under cutoff regularization, vanishes under dimensional regularization, and has no BISH-computable empirical content. (II) The "naturalness" argument is a Bayesian prior, not a mathematical derivation; it resides outside the BISH/LPO deductive hierarchy. (III) The genuine constraint is the 55-decimal-place cancellation between the bare cosmological constant and the electroweak/QCD vacuum condensates. Computing the exact interacting condensates requires the thermodynamic limit (Fekete's lemma, Paper 29), showing that LPO suffices for the fine-tuning equation. The cosmological constant problem introduces no new logical resources. The BISH+LPO ceiling holds. The entire formalization (10 Lean 4/Mathlib4 modules, ~830 lines) compiles with zero `sorry`, zero warnings.

1 Introduction

The cosmological constant problem is widely described as the worst prediction in the history of physics. Quantum field theory allegedly predicts a vacuum energy density of order $M_{\text{Planck}}^4 \approx 10^{71}$ GeV^4, while the observed value is $\rho_\Lambda \approx 10^{-47}$ GeV^4 —a discrepancy of 120 orders of magnitude [1]. The problem traces to Zel'dovich's [16] identification of quantum vacuum energy as a gravitational source, was crystallized in Weinberg's landmark review [1], and remains unsolved despite decades of sustained effort (see [23–25] for comprehensive reviews). On the observational side, the discovery of cosmic acceleration via Type Ia supernovae [17, 18], precision constraints from Planck [19], and recent DESI baryon acoustic oscillation measurements [20] have confirmed that a small positive cosmological constant—or something closely mimicking it—is required by the data.

This paper subjects the cosmological constant problem to the axiom calibration framework of constructive reverse mathematics (CRM), developed across Papers 1–41 of this programme. The central question is not "how do we solve the problem?" but rather "what is the logical

*New York University. AI-assisted formalization; see §17 for methodology.

structure of each component?" At what level of the constructive hierarchy does each piece of the alleged prediction live?

Constructive reverse mathematics. Constructive mathematics, as developed by Bishop [6, 7, 21], takes intuitionistic logic as its foundation: proofs must provide explicit computational witnesses rather than relying on the law of excluded middle (LEM) or unrestricted choice. Constructive reverse mathematics (CRM) [22] asks a complementary question: for each theorem in mathematics, what is the *minimal* logical principle beyond the constructive base BISH that suffices to prove it?

The answer organizes into a hierarchy called the *omniscience spine* [8]:

$$\text{BISH} \;<\; \text{LLPO} \;<\; \text{WLPO} \;<\; \text{LPO} \;<\; \text{LEM}.$$

Each level corresponds to a specific computational capability. BISH (Bishop-style constructive mathematics) suffices for any finite computation: lattice sums, Born probabilities, Heisenberg uncertainty bounds. LLPO (the lesser limited principle of omniscience) adds the ability to assert a disjunction without a constructive witness for either disjunct—it governs Bell's theorem, intermediate value root-finding, and WKB turning-point decisions. WLPO (the weak limited principle) decides whether an infinite binary sequence is "all zeros" without producing a counterexample; it governs the bidual gap, singular states, and phase classification. LPO (the limited principle of omniscience) provides full binary decidability: either the sequence is all zeros, or we can exhibit a nonzero term. It governs every thermodynamic limit in the programme via Fekete's subadditive lemma [5, 11]. The Fan Theorem (FT), governing compactness arguments, is independent of the entire spine. Paper 10 [9] assembles a comprehensive calibration table of approximately 50 entries across 11 physical domains; Paper 12 [10] provides a historical narrative showing how these non-constructive commitments entered physics through Weierstrass's analysis and Boltzmann's statistical mechanics. Paper 40 [14] consolidates Papers 1–41 into a monograph-length summary, defending the BISH + LPO characterization and providing the programme-wide context for the present paper.

Axiom calibration. The programme's methodology proceeds as follows. For each physical theorem, we identify (i) the *bridge axioms* that encode the physical content (e.g., "the lattice energy is subadditive," "the mode sum is monotone"), and (ii) the minimal CRM principle the proof requires. The result is a calibration: a mapping from physical content to logical cost. A quantity is *scaffolding* if its value depends on a mathematical choice (regularization scheme, basis, gauge) that has no empirical counterpart. All formalizations are machine-checked in Lean 4 with Mathlib [9].

The central finding across 41 papers: most empirical predictions in physics—finite-size bounds, Born probabilities, uncertainty relations, Casimir energy differences, perturbative cross-sections—are BISH-computable. Thermodynamic limits (free energy densities, exact condensates, geodesic completeness) universally cost LPO, via the equivalence of Fekete's subadditive lemma with LPO [11]. The BISH + LPO ceiling has held across all domains examined.

The cosmological constant problem in context. The cosmological constant problem is often presented as a single puzzle, but it conflates at least three logically distinct issues [1, 23]: (i) the *old cosmological constant problem* (why is ρ_{vac} not of order M_{Planck}^4?), (ii) the *coincidence problem* (why is ρ_Λ comparable to the matter density today?), and (iii) the *naturalness problem* (why does Λ not receive large quantum corrections?).

The standard narrative relies on a calculation that is more subtle than often acknowledged. Martin [25] demonstrates pedagogically that the 10^{120} figure arises specifically from hard-cutoff regularization; dimensional regularization—which preserves gauge and Lorentz invariance—removes power-law divergences entirely and produces a qualitatively different (and much smaller) result. The Hollands–Wald framework [2, 3] proves rigorously that in quantum field theory on curved spacetime, the renormalized stress-energy tensor contains a free parameter c_1 that plays the role of the cosmological constant. This parameter cannot be predicted by QFT; it must be fixed by observation.

The assumption that Λ "should" be of order M^4_{Planck} derives from 't Hooft's [26] concept of technical naturalness: a parameter is natural only if setting it to zero increases the symmetry of the theory. Burgess [27] develops this into an effective field theory framework for the CC problem. However, the naturalness criterion has faced mounting criticism. Bianchi and Rovelli [28] argue that Λ should simply be treated as a fundamental constant of nature, like G or c, requiring no more explanation than they do. Hossenfelder [29] argues that fine-tuning arguments lack a well-defined probability distribution over parameter space and are therefore not scientifically productive. Giudice [30] diagnoses a "post-naturalness era" in which the LHC results have undermined the predictive utility of naturalness reasoning.

Nobbenhuis [31] categorizes proposed solutions into five classes: symmetry mechanisms, back-reaction, modifications of general relativity, statistical/anthropic approaches, and vacuum-energy adjustment mechanisms. The string landscape [32, 33] provides a paradigmatic example of the last category. The present paper does not propose a new approach in any of these categories. Instead, it asks a different question: what is the *logical structure* of each component of the problem?

Contribution of this paper. We apply the axiom calibration framework to decompose the cosmological constant problem into three logically distinct claims with different constructive status. This analysis is orthogonal to the approaches catalogued by Nobbenhuis [31]: it does not propose a mechanism, but identifies the logical prerequisites.

- **Claim I (Dissolved):** The 10^{120} is a regulator-dependent artifact. It is scaffolding, not a prediction.
- **Claim II (Reclassified):** Naturalness is a Bayesian prior, not a theorem. It resides outside the constructive hierarchy.
- **Claim III (Identified,** LPO**):** The 55-digit fine-tuning is real. It is an arithmetic relation between LPO-computable reals—the same logical level as every thermodynamic limit in the programme.

Claim I: The 10^{120} UV discrepancy $\longrightarrow$ **DISSOLVED**
Regulator-dependent artifact. Vanishes under dim. reg. *Thms 4.1, 5.1* — BISH

Claim II: Naturalness ($\Lambda \sim M^4_{\text{Planck}}$) $\longrightarrow$ **RECLASSIFIED**
Bayesian prior, not mathematical derivation. *Thm 6.1* — BISH

Claim III: 55-digit fine-tuning $\longrightarrow$ **IDENTIFIED**
$\Lambda_{\text{obs}} = \Lambda_{\text{bare}} + 8\pi G(\rho_H^{\text{exact}} + \rho_{\text{QCD}}^{\text{exact}})$. *Thms 8.1, 10.1* — LPO

Figure 1: The cosmological constant problem decomposed. The 10^{120} discrepancy is dissolved as scaffolding. Naturalness is reclassified as non-mathematical. The genuine constraint is a 55-digit fine-tuning at LPO.

Formalization overview. The Lean 4 bundle consists of 10 modules totalling ∼830 lines. Seven theorems establish the decomposition. The assembly theorem `cc_calibration` combines all seven parts into a single conjunction. The master theorem `cc_master` re-exports the assembly. The axiom audit (`#print axioms cc_master`) confirms a clean profile: 11 physics bridge axioms, 1 CRM axiom (`bmc_from_subadditive`), and the standard Lean infrastructure triple (`propext`, `Classical.choice`, `Quot.sound`).

2 The Three Claims

2.1 Claim I: The 10^{120} Is Scaffolding

For a free scalar field of mass m on a finite lattice, the vacuum energy $E_{\text{vac}} = \frac{1}{2}\sum_k \omega_k$ is a finite sum of algebraic expressions. BISH. In the continuum limit, the sum becomes a quartically divergent integral. It does not converge; it does not define a real number at any level.

Different regularization schemes extract different finite numbers:

- **Cutoff:** $\rho \sim \Lambda_{\text{UV}}^4/(16\pi^2)$. Setting $\Lambda_{\text{UV}} = M_{\text{Planck}}$ gives $\rho \sim 10^{71}$ GeV^4.
- **Dimensional:** Power-law divergences vanish. $\rho \sim m^4 \ln(m^2/\mu^2) \sim (100 \text{ GeV})^4$.
- **ζ-function:** Agrees with dimensional regularization.

The "prediction" changes with the scaffolding. A regulator-dependent quantity has no empirical content.

SCAFFOLDING		PHYSICAL
Absolute vacuum energy	diagnostic	**Energy differences**
Cutoff: $\sim M_{\text{Planck}}^4$		Casimir: $-\pi^2/(240\, d^4)$
Dim. reg.: $\sim m^4 \ln$		All regulators agree
ζ-function: $\sim m^4 \ln$		Experimentally verified
Regulator-dependent		BISH-**computable**

Figure 2: The scaffolding diagnostic applied to vacuum energy. Absolute vacuum energies (left) depend on the regularization scheme and carry no empirical content. Energy *differences* (right) are scheme-independent, BISH-computable, and experimentally verified (Casimir effect).

Remark 2.1 (Scope of dissolution)**.** The dissolution applies to the *magnitude* of the discrepancy (10^{120}), not to the *existence* of individual $m^4/(16\pi^2)$ contributions from massive particles. Each species contributes a quartic term that is finite and scheme-independent in form; what is scheme-dependent is the *sum* over all modes including the UV region. The framework dissolves the UV catastrophe narrative, not the perturbative physics of individual particle thresholds.

2.2 Claim II: Naturalness Is Not Mathematics

The Hollands–Wald axioms [2, 3] prove that the renormalized stress tensor $\langle T_{\mu\nu}\rangle_{\text{ren}}$ is determined up to free local geometric coefficients: c_1 is the cosmological constant. It is a *free parameter*, on the same footing as particle masses and coupling constants. The "naturalness" expectation that $c_1 \sim M_{\text{Planck}}^4$ is a Bayesian prior—a claim about expected magnitudes, not a derivation.

2.3 Claim III: The 55-Digit Fine-Tuning

After dissolving Claim I and reclassifying Claim II, the genuine constraint is:

$$\Lambda_{\text{obs}} = \Lambda_{\text{bare}} + 8\pi G\big(\rho_{\text{Higgs}}^{\text{exact}} + \rho_{\text{QCD}}^{\text{exact}}\big).$$

The tree-level condensates are BISH (algebraic expressions of measured parameters: $\rho_{\text{Higgs}} \approx -\mu^4/(4\lambda)$, $\rho_{\text{QCD}} \approx -\langle\bar{q}q\rangle m_q$). The exact interacting condensates—including all loop corrections, non-perturbative effects, and vacuum fluctuations—require the thermodynamic limit (Fekete, LPO; Paper 29).

Remark 2.2 (Coincidence problem)**.** This paper addresses only the "old" cosmological constant problem (why is Λ_{obs} so small?) and not the coincidence problem (why is $\rho_\Lambda \sim \rho_{\text{matter}}$ today?). The coincidence problem involves cosmological dynamics and is out of scope for the present axiom calibration.

The fine-tuning equation is an LPO equality. The "55-digit" figure is computed as follows: the dominant condensate is the Higgs at $|\rho_H| \sim (100\text{ GeV})^4 = 10^8\text{ GeV}^4$, while $\Lambda_{\text{obs}}/8\pi G \sim 10^{-47}\text{ GeV}^4$. The cancellation ratio is $10^8/10^{-47} = 10^{55}$, requiring Λ_{bare} and $8\pi G\rho_{\text{exact}}$ to agree to 55 decimal places.

3 Core Definitions and Bridge Axioms

The formalization begins with `Defs.lean`, which establishes the type-theoretic infrastructure. We highlight the key design decisions.

Definition 3.1 (Regularization Scheme)**.** A regularization scheme is modeled as an inductive type with three constructors: hard momentum cutoff (carrying $\Lambda_{\text{UV}} > 0$), dimensional regularization (carrying $\mu > 0$), and ζ-function regularization (no parameters).

```
-- Regularization schemes
inductive RegScheme : Type where
  | cutoff (L_UV : R) (hL : 0 < L_UV)
  | dimreg (mu : R) (hmu : 0 < mu)
  | zeta

-- Wald ambiguity structure
structure WaldAmbiguity where
  c1 : R     -- free parameter
  condensate_sum : R

def effective_Lambda (w : WaldAmbiguity) : R :=
  w.c1 + w.condensate_sum

-- Bridge axioms (physics -> computability)
axiom regularized_vacuum_energy : RegScheme -> R
axiom mode_sum_partial (m : R) : N -> R
axiom mode_sum_unbounded (m : R) (hm : 0 <= m) :
    forall M : R, exists N : N, mode_sum_partial m N > M
axiom lattice_energy_subadditive :
    forall m n : N, lattice_vacuum_energy (m + n) <=
              lattice_vacuum_energy m + lattice_vacuum_energy n
axiom casimir_cauchy_modulus (d : R) (hd : 0 < d) :
    forall e : R, 0 < e ->
      exists (N : N) (approx : R),
        N > 0 /\ |approx - (-Real.pi ^ 2
          / (240 * d ^ 4))| < e
axiom picard_lindelof_lambda
    (mu_min mu_max : R) (h_min : 0 < mu_min)
    (h_range : mu_min < mu_max) (L_init : R) :
    exists L_sol : R -> R,
      L_sol mu_min = L_init /\
      forall mu : R, mu_min <= mu -> mu <= mu_max ->
        forall e : R, 0 < e -> exists d : R, 0 < d /\
          forall mu' : R, |mu' - mu| < d ->
            |L_sol mu' - L_sol mu| < e
```

Listing 1: Core types: `Defs.lean` (excerpts)

Remark 3.2 (Bridge axiom philosophy)**.** Each axiom encodes a single, well-established physical fact: unboundedness of the mode sum, subadditivity of the ground-state energy, exponential decay of the Casimir remainder integrand, Lipschitz regularity of the RG beta function. These are unverified premises in the formal sense: the Lean type-checker guarantees that the theorems follow *from* the axioms, but does not verify the axioms themselves. The bridge axioms are the interface between the mathematical framework and the physical content.

Remark 3.3 (Constructive status of bridge axioms)**.** The bridge axioms themselves have varying constructive status.

- *Monotonicity and non-negativity* of mode sums: BISH-provable from the definition of $\omega_k = \sqrt{k^2 + m^2}$.
- *Unboundedness* of the mode sum: BISH-provable by explicit lower bound $E_N \geq N \cdot m/2$ for N large.
- *Subadditivity* of lattice ground-state energy: BISH-provable for short-range interactions (the boundary contribution is bounded by a surface-area term). For QCD, subadditivity relies on confinement/cluster decomposition; for QED, long-range Coulomb interactions require Debye screening. See Remark 8.2 for details.
- *Casimir modulus*: BISH-provable from the exponential decay of the Abel–Plana remainder $\sim e^{-2\pi t}$.
- *Picard–Lindelöf*: a BISH theorem (constructive contraction mapping); declared as axiom to avoid re-proving the general ODE theorem within this project.
- *Regulator dependence*: follows from explicit calculation; the cutoff result ($\sim \Lambda_{\mathrm{UV}}^4$) differs from the dim. reg. result ($\sim m^4 \log$) by inspection.

Thus most bridge axioms are in principle BISH-provable from more primitive physical assumptions. They are declared as axioms for modularity, not because their constructive status is in doubt.

4 Theorem 1: Vacuum Energy Diverges

Theorem 4.1 (Vacuum Energy Divergence)**.** *The unregularized vacuum energy mode sum $E_N = \frac{1}{2}\sum_{|k|\leq N}\sqrt{k^2+m^2}$ is unbounded. No limit $L \in \mathbb{R}$ exists:*

$$\neg\, \exists L \in \mathbb{R},\ \forall \varepsilon > 0,\ \exists N_0,\ \forall N \geq N_0,\ |E_N - L| < \varepsilon.$$

Proof. Suppose for contradiction that $E_N \to L$. For $\varepsilon = 1$, convergence gives an index N_0 with $|E_N - L| < 1$ for all $N \geq N_0$, i.e., $L - 1 < E_N < L + 1$. By the unboundedness bridge axiom, there exists N_1 with $E_{N_1} > L+1$. Since the mode sum is monotonically increasing (more modes means more positive-definite energy), we have $E_{\max(N_0,N_1)} \geq E_{N_1} > L + 1$. But simultaneously $|E_{\max(N_0,N_1)} - L| < 1$, which forces $E_{\max(N_0,N_1)} < L + 1$. Contradiction.

BISH: the proof uses only unboundedness, monotonicity, and the triangle inequality. No omniscience principles. □

Remark 4.2. The "continuum vacuum energy" is not a real number—not at BISH, not at LPO, not anywhere in the constructive hierarchy. A divergent series has no BMC limit. The 10^{120} narrative begins with a quantity that does not exist as a mathematical object.

```
theorem vacuum_energy_divergent (m : R) (hm : 0 <= m) :
    not exists L : R, forall e : R, 0 < e ->
      exists N0 : N, forall N : N, N0 <= N ->
        |mode_sum_partial m N - L| < e := by
  intro ⟨L, hconv⟩
  obtain ⟨N1, hN1⟩ := mode_sum_unbounded m hm (L + 1)
  obtain ⟨N0, hN0⟩ := hconv 1 one_pos
  have hN := hN0 (max N0 N1) (le_max_left N0 N1)
  have hbig : mode_sum_partial m (max N0 N1) > L + 1 :=
    lt_of_lt_of_le hN1
```

```
    (mode_sum_mono m hm (le_max_right N0 N1))
  linarith [abs_lt.mp hN]
```

Listing 2: Vacuum energy divergence: `VacuumDivergence.lean`

Corollary 4.3 (Not Cauchy)**.** *The vacuum energy mode sum is not Cauchy: there is no modulus of convergence. If the partial sums were Cauchy, they would be bounded (take $\varepsilon = 1$; all terms beyond N_0 lie within 1 of E_{N_0}). But the mode sum is unbounded. Contradiction.*

```
theorem vacuum_energy_not_cauchy (m : R) (hm : 0 <= m) :
    not forall e : R, 0 < e ->
      exists N0 : N, forall M N : N, N0 <= M -> N0 <= N ->
        |mode_sum_partial m M - mode_sum_partial m N| < e := by
  intro hcauchy
  obtain ⟨N0, hN0⟩ := hcauchy 1 one_pos
  have hbdd : forall N : N, N0 <= N ->
      mode_sum_partial m N < mode_sum_partial m N0 + 1 := by
    intro N hN
    have := hN0 N0 N le_rfl hN
    linarith [abs_lt.mp this]
  obtain ⟨N1, hN1⟩ :=
    mode_sum_unbounded m hm (mode_sum_partial m N0 + 1)
  have := hbdd (max N0 N1) (le_max_left _ _)
  have : mode_sum_partial m (max N0 N1) >
      mode_sum_partial m N0 + 1 :=
    lt_of_lt_of_le hN1
      (mode_sum_mono m hm (le_max_right _ _))
  linarith
```

Listing 3: Not Cauchy: `VacuumDivergence.lean`

5 Theorem 2: Regulator Dependence

Theorem 5.1 (Regulator Dependence)**.** *There exist two valid regularization schemes r_1, r_2 that yield different absolute vacuum energies:*

$$\exists\, r_1, r_2 : \mathrm{RegScheme}, \quad \rho(r_1) \neq \rho(r_2).$$

Proof. Cutoff regularization with $\Lambda_{\mathrm{UV}} = M_{\mathrm{Planck}}$ yields $\rho > 0$ (quartic in Λ_{UV}). Dimensional regularization with scale μ yields a different value ($\sim m^4 \ln(m^2/\mu^2)$, generically different from the quartic form). The bridge axiom `dimreg_value_different` asserts that for some valid parameters, the two values differ. The proof constructs the two `RegScheme` terms and applies the axiom.

BISH: the proof is a direct existential witness. No limits. □

Remark 5.2. This is the formal expression of Claim I. A quantity that depends on the scaffolding choice cannot be a physical prediction. The calibration framework's criterion: empirical content must be invariant under change of scaffolding. The vacuum energy fails this criterion.

```
theorem prediction_regulator_dependent :
    exists (r1 r2 : RegScheme),
      regularized_vacuum_energy r1 !=
      regularized_vacuum_energy r2 := by
  obtain ⟨L_UV, hL, mu, hmu, hne⟩ := dimreg_value_different
  exact ⟨RegScheme.cutoff L_UV hL,
```

```
          RegScheme.dimreg mu hmu, hne⟩

-- Corollary: no invariant vacuum energy exists
theorem no_regulator_invariant_vacuum_energy :
    not exists (rho : R),
      forall r : RegScheme,
        regularized_vacuum_energy r = rho := by
  intro ⟨rho, hinv⟩
  obtain ⟨L_UV, hL, mu, hmu, hne⟩ := dimreg_value_different
  exact hne (by rw [hinv, hinv])
```

Listing 4: Regulator dependence: `RegulatorDependence.lean`

6 Theorem 3: Wald Ambiguity (BISH)

Theorem 6.1 (Wald Ambiguity)**.** *For any target $\Lambda_{\text{obs}} \in \mathbb{R}$ and any condensate contribution $C \in \mathbb{R}$, there exists a valid Wald ambiguity parameter c_1 such that $\Lambda_{\text{eff}} = c_1 + C = \Lambda_{\text{obs}}$.*

Proof. Set $c_1 = \Lambda_{\text{obs}} - C$. Then $\Lambda_{\text{eff}} = c_1 + C = (\Lambda_{\text{obs}} - C) + C = \Lambda_{\text{obs}}$. The proof is subtraction followed by `ring`.

BISH: no limits, no compactness, no choice principles. This is the simplest proof in the entire formalization—pure arithmetic. The Hollands–Wald theorem ensures that c_1 is a legitimate free parameter of the renormalized theory, not an ad hoc patch. □

Remark 6.2. This theorem formalizes the Hollands–Wald result: QFT in curved spacetime *cannot predict* the cosmological constant. The coefficient c_1 is a free parameter, settable by arithmetic. For *any* real number r, there exists a valid theory with $\Lambda_{\text{eff}} = r$. The "prediction" of a specific value is a confusion between the formalism (which leaves c_1 undetermined) and the scaffolding (which assigns an arbitrary value to c_1 via the regulator).

```
theorem lambda_free_parameter
    (L_obs condensate_sum : R) :
    exists w : WaldAmbiguity,
      w.condensate_sum = condensate_sum /\
      effective_Lambda w = L_obs := by
  exact ⟨{c1 := L_obs - condensate_sum, condensate_sum},
    rfl, by unfold effective_Lambda; ring⟩

-- QFT cannot predict Lambda: for ANY r, a valid theory exists
theorem qft_cannot_predict_lambda
    (r : R) (condensate_sum : R) :
    exists w : WaldAmbiguity,
      w.condensate_sum = condensate_sum /\
      effective_Lambda w = r :=
  lambda_free_parameter r condensate_sum

-- The space of valid theories is affine in c1
theorem wald_ambiguity_affine (condensate_sum : R) :
    forall c : R,
      effective_Lambda {c1 := c, condensate_sum} =
        c + condensate_sum := by
  intro c; rfl
```

Listing 5: Wald ambiguity: `WaldAmbiguity.lean`

7 Theorem 4: Casimir Energy (BISH)

Theorem 7.1 (Casimir Energy Is BISH)**.** *The Casimir energy difference between conducting plates of separation $d > 0$ converges to $-\pi^2/(240\,d^4)$ with a computable Cauchy modulus. For every $\varepsilon > 0$, there exist $N \in \mathbb{N}$ and an approximation $a \in \mathbb{R}$ with $|a - E| < \varepsilon$.*

Proof. The Casimir force is an energy *difference*: $E_{\text{plates}} - E_{\text{free}}$. When one computes the difference of the mode sums for the plate and free-space configurations, the divergent (regulator-dependent) terms cancel algebraically. After Abel–Plana or Euler–Maclaurin summation, the finite remainder involves an integrand with exponential decay $\sim (e^{2\pi t} - 1)^{-1}$. This exponential decay rate provides an explicit, computable error bound: truncating the quadrature at N terms introduces an error bounded by $C \cdot e^{-2\pi N}$. For any target ε, choosing $N > \ln(C/\varepsilon)/(2\pi)$ suffices.

BISH: finite computation with guaranteed error bounds. No LPO required. The key observation: QFT knows the difference between absolute energies (regulator-dependent, no physical content) and energy *differences* (BISH, experimentally verified to 1% accuracy [4]). $\square$

Remark 7.2 (The scaffolding diagnostic)**.** The Casimir effect is the paradigm case for the framework's scaffolding diagnostic. Absolute vacuum energies are regulator-dependent (no physical content). Energy *differences* are BISH and match experiment. The formalization captures this dichotomy: `casimir_energy` is defined as the finite quantity $-\pi^2/(240d^4)$, and `casimir_is_energy_difference` proves it is negative (attractive force between the plates).

```
def casimir_energy (d : R) : R :=
  -Real.pi ^ 2 / (240 * d ^ 4)

theorem casimir_bish (d : R) (hd : 0 < d) :
    exists (E : R),
      E = casimir_energy d /\
      forall e : R, 0 < e ->
        exists (N : N) (approx : R),
          N > 0 /\ |approx - E| < e := by
  refine ⟨casimir_energy d, rfl, ?_⟩
  intro e he
  obtain ⟨N, approx, hN, h⟩ :=
    casimir_cauchy_modulus d hd e he
  exact ⟨N, approx, hN,
    by unfold casimir_energy; exact h⟩

-- Casimir energy is negative (attractive force)
theorem casimir_is_energy_difference
    (d : R) (hd : 0 < d) :
    casimir_energy d < 0 := by
  unfold casimir_energy
  have h240d : (0 : R) < 240 * d ^ 4 := by positivity
  have hpi : (0 : R) < Real.pi ^ 2 := by positivity
  simp only [neg_div]
  exact neg_neg_of_pos (div_pos hpi h240d)
```

Listing 6: Casimir energy: `CasimirBISH.lean`

8 Theorem 5: Condensate (LPO)

Theorem 8.1 (Condensate Is LPO)**.** *Assuming* LPO*, the exact interacting vacuum energy density exists:*

$$\text{LPO} \longrightarrow \exists\, \rho_{\text{exact}} \in \mathbb{R},\ \forall \varepsilon > 0,\ \exists N_0,\ \forall n \geq N_0,\ n > 0 \Rightarrow \left| \frac{E(n)}{n} - \rho_{\text{exact}} \right| < \varepsilon.$$

Proof. The ground-state energy $E(L)$ on a lattice of volume L satisfies two bridge axioms:

1. **Subadditivity:** $E(m+n) \leq E(m) + E(n)$ for all $m, n \in \mathbb{N}$. (Physical origin: translation invariance of the Hamiltonian; the boundary interaction between two subsystems is bounded.)
2. **Bounded below:** $\exists C \in \mathbb{R}$, $\forall n > 0$, $C \leq E(n)/n$. (Physical origin: the energy density cannot diverge to $-\infty$; stability of the vacuum.)

Remark 8.2 (Justification of subadditivity)**.** Subadditivity holds rigorously for short-range lattice systems: if the interaction decays exponentially, the boundary energy between two subsystems of sizes m and n is bounded by a surface-area term $O(L^{d-1})$, which is sub-extensive and yields $E(m+n) \leq E(m) + E(n) + O((m+n)^{(d-1)/d})$—sufficient for Fekete's lemma to apply. For QCD, subadditivity relies on confinement: color-singlet states have exponentially decaying correlations (cluster decomposition), making the effective interaction short-range. For QED, Coulomb interactions are long-range, but Debye screening in the vacuum (from virtual e^+e^- pairs) renders the effective interaction short-range at distances beyond the Compton wavelength $\sim 1/m_e$. A rigorous proof of subadditivity for the full Standard Model lattice Hamiltonian remains an open mathematical problem; the bridge axiom encodes the standard physics expectation.

By Fekete's subadditive lemma, the limit $\ell = \inf_{n\geq 1} E(n)/n$ exists and the sequence $E(n)/n$ converges to ℓ. By Paper 29, Fekete's lemma is equivalent to LPO (via Bounded Monotone Convergence). The CRM axiom `bmc_from_subadditive` captures this equivalence.

The proof is a single function application: supply the subadditivity and boundedness bridge axioms, plus the LPO hypothesis, to obtain the convergent density. □

Remark 8.3 (Tree level vs. exact)**.** The tree-level Higgs condensate $\rho_{\text{Higgs}} \approx -\mu^4/(4\lambda) \approx -(100\text{ GeV})^4$ and the tree-level QCD condensate $\rho_{\text{QCD}} \approx -\langle \bar{q}q \rangle m_q \approx -(200\text{ MeV})^3 \cdot (\text{few MeV})$ are both BISH: algebraic expressions of measured parameters, requiring no limits. The *exact* interacting values—including all loop corrections and non-perturbative effects—require the thermodynamic limit, hence LPO. This is the same BISH → LPO step found throughout the programme (Ising model in Paper 8, lattice gauge theory in Paper 33, etc.).

```
theorem condensate_lpo :
    LPO -> exists rho_exact : R,
      forall e : R, 0 < e ->
        exists N0 : N, forall n : N,
          n >= N0 -> 0 < n ->
            |lattice_vacuum_energy n / (n : R)
              - rho_exact| < e := by
  intro lpo
  exact bmc_from_subadditive
    lattice_vacuum_energy
    lattice_energy_subadditive
    lattice_energy_bdd_below lpo

-- Tree-level condensates are BISH (no LPO needed)
theorem higgs_tree_bish :
    exists rho_H : R,
      rho_H = higgs_tree_level /\ rho_H < 0 :=
  ⟨higgs_tree_level, rfl, higgs_tree_level_neg⟩

theorem qcd_tree_bish :
    exists rho_QCD : R,
      rho_QCD = qcd_tree_level /\ rho_QCD < 0 :=
  ⟨qcd_tree_level, rfl, qcd_tree_level_neg⟩
```

Listing 7: Condensate is LPO: `CondensateLPO.lean`

9 Theorem 6: RG Running (BISH)

Theorem 9.1 (RG Running Is BISH)**.** *The perturbative RG running of Λ is* BISH*-computable. For any initial $\Lambda(\mu_0)$ and target scale μ in the perturbative regime ($\mu > \Lambda_{\mathrm{QCD}}$), the running $\Lambda(\mu)$ is a* BISH*-computable real.*

Proof. The RG equation $\mu\, d\Lambda/d\mu = \beta_\Lambda(\mu)$ is a first-order ODE. The beta function β_Λ is a finite sum of contributions from the Standard Model spectrum—algebraic functions of particle masses and couplings:

$$\beta_\Lambda(\mu) = \frac{1}{16\pi^2} \sum_{i \in \mathrm{SM}} (-1)^{F_i}\, n_i\, m_i^4\, f\left(\frac{m_i^2}{\mu^2}\right).$$

The right-hand side is Lipschitz on any bounded interval $[\mu_{\min}, \mu_{\max}]$ (away from particle thresholds where $\mu = m_i$).

By the Picard–Lindelöf theorem—which is constructive for Lipschitz ODEs via the Banach contraction mapping principle—the solution exists, is unique, and the Picard iterates provide a computable approximation scheme with explicit error bounds (the Euler method with step size h has error $O(h)$).

LPO enters only below the QCD scale $\mu \lesssim \Lambda_{\mathrm{QCD}} \approx 200$ MeV, where the non-perturbative QCD condensate contributes. This non-perturbative contribution requires the thermodynamic limit (Theorem 8.1). Above Λ_{QCD}, everything is perturbative and BISH. $\square$

```
theorem rg_running_bish (mu0 mu_target : R)
    (h0 : 0 < mu0) (h_range : mu0 < mu_target)
    (L_init : R) :
    exists L_sol : R -> R,
      L_sol mu0 = L_init /\
      forall mu : R, mu0 <= mu -> mu <= mu_target ->
        forall e : R, 0 < e -> exists d : R, 0 < d /\
          forall mu' : R, |mu' - mu| < d ->
            |L_sol mu' - L_sol mu| < e := by
  exact picard_lindelof_lambda mu0 mu_target
    h0 h_range L_init

-- LPO boundary: perturbative is BISH,
-- non-perturbative QCD requires LPO
theorem rg_lpo_boundary :
    (forall (mu0 mu : R) (_h0 : 0 < mu0)
       (_hr : mu0 < mu) (L_init : R),
      exists L_sol : R -> R,
        L_sol mu0 = L_init) /\
    (LPO -> exists rho_QCD_exact : R,
      forall e : R, 0 < e ->
        exists N0 : N, forall n : N,
          n >= N0 -> 0 < n ->
            |lattice_vacuum_energy n / (n : R)
              - rho_QCD_exact| < e) := by
  constructor
  . intro mu0 mu h0 hr L_init
    obtain ⟨sol, hsol, _⟩ :=
      picard_lindelof_lambda mu0 mu h0 hr L_init
```

```
    exact ⟨sol, hsol⟩
  . exact bmc_from_subadditive
      lattice_vacuum_energy
      lattice_energy_subadditive
      lattice_energy_bdd_below
```

Listing 8: RG running: `RGRunningBISH.lean`

10 Theorem 7: Fine-Tuning (LPO)

Theorem 10.1 (Fine-Tuning Is LPO)**.** *The fine-tuning equation* $\Lambda_{\mathrm{obs}} = \Lambda_{\mathrm{bare}} + 8\pi G\,\rho_{\mathrm{exact}}$ *is an equality between* LPO*-computable reals.*

Proof. Given LPO:

1. The exact condensate ρ_{exact} exists (Theorem 8.1, via Fekete/LPO).
2. The gravitational coupling $8\pi G$ is a measured constant (BISH).
3. The product $8\pi G \cdot \rho_{\mathrm{exact}}$ is an arithmetic operation on an LPO-computable real (LPO).
4. Λ_{bare} is a free parameter (Theorem 6.1, BISH). Set $\Lambda_{\mathrm{bare}} = \Lambda_{\mathrm{obs}} - 8\pi G\,\rho_{\mathrm{exact}}$.

The equation holds by arithmetic (`ring`). The LPO cost is inherited entirely from the Fekete limit in step 1, not from anything intrinsic to the equation or its structure. □

Remark 10.2. The 55-digit cancellation is *real* but *logically mundane*. It is an arithmetic relation between LPO-computable reals—the same logical level as the Ising model phase transition temperature (Paper 8), the QCD confinement scale (Paper 33), and every other thermodynamic limit in the programme. The framework does *not* explain why the cancellation occurs. It identifies the question "why is Λ_{obs} so small?" as a question about initial conditions, not about the logical structure of QFT.

```
theorem fine_tuning_lpo :
    LPO ->
    exists (rho_exact L_bare : R),
      Lambda_obs = L_bare + eight_pi_G * rho_exact := by
  intro lpo
  obtain ⟨rho, _h_rho⟩ := condensate_lpo lpo
  -- Wald ambiguity: set L_bare = L_obs - 8piG * rho
  exact ⟨rho, Lambda_obs - eight_pi_G * rho,
    by ring⟩

-- Tree-level fine-tuning is already BISH
theorem fine_tuning_tree_bish :
    exists L_bare : R,
      Lambda_obs = L_bare + eight_pi_G *
        (higgs_tree_level + qcd_tree_level) := by
  exact ⟨Lambda_obs - eight_pi_G *
    (higgs_tree_level + qcd_tree_level), by ring⟩

-- The LPO cost is entirely from the thermodynamic limit
theorem lpo_cost_is_thermodynamic :
    (exists L_bare : R,
      Lambda_obs = L_bare + eight_pi_G *
        (higgs_tree_level + qcd_tree_level)) /\
    (LPO -> exists rho_exact : R,
      forall e : R, 0 < e ->
        exists N0 : N, forall n : N,
          n >= N0 -> 0 < n ->
```

```
            |lattice_vacuum_energy n / (n : R)
              - rho_exact| < e) := by
  exact ⟨fine_tuning_tree_bish, condensate_lpo⟩
```

Listing 9: Fine-tuning: FineTuningLPO.lean

11 Assembly and Master Theorem

The calibration table and assembly theorem are defined in CalibrationTable.lean. The assembly theorem cc_calibration combines all seven parts into a single conjunction, which the master theorem cc_master re-exports.

```
structure CalibrationEntry where
  name : String
  constructive_status : String
  physical_content : String
  deriving Repr

def calibration_table : List CalibrationEntry := [
  ⟨"Mode sum on finite lattice",
    "BISH", "Yes (lattice QFT)"⟩,
  ⟨"Continuum limit of mode sum",
    "Divergent (no real number)", "None"⟩,
  ⟨"Casimir energy difference",
    "BISH", "Yes (matches experiment)"⟩,
  ⟨"Exact interacting condensate",
    "LPO (Fekete)", "Yes (enters Einstein eq.)"⟩,
  ⟨"Fine-tuning equation (exact)",
    "LPO", "Yes"⟩,
  -- ... (12 rows total)
]

theorem cc_calibration :
    -- Part 1: Vacuum energy diverges
    (...) /\
    -- Part 2: Regulator-dependent
    (...) /\
    -- Part 3: Lambda is free (BISH)
    (...) /\
    -- Part 4: Casimir is BISH
    (...) /\
    -- Part 5: Condensate is LPO
    (...) /\
    -- Part 6: RG running is BISH
    (...) /\
    -- Part 7: Fine-tuning is LPO
    (...) :=
  ⟨vacuum_energy_divergent,
   prediction_regulator_dependent,
   fun L_obs condensate => ⟨...⟩,
   fun d _hd => ⟨casimir_energy d, rfl⟩,
   condensate_lpo,
   fun mu0 mu h0 hr L_init => by
     obtain ⟨sol, hsol, _⟩ :=
       picard_lindelof_lambda mu0 mu h0 hr L_init
     exact ⟨sol, hsol⟩,
   fine_tuning_lpo⟩
```

Listing 10: Assembly: `CalibrationTable.lean` (excerpt)

```
theorem cc_master :
    -- (7-part conjunction, same type as cc_calibration)
    ... :=
  cc_calibration

#print axioms cc_master
-- 'cc_master' depends on axioms:
-- [Lambda_obs, bmc_from_subadditive,
--   dimreg_value_different, eight_pi_G,
--   lattice_energy_bdd_below,
--   lattice_energy_subadditive,
--   lattice_vacuum_energy, mode_sum_mono,
--   mode_sum_partial, mode_sum_unbounded,
--   picard_lindelof_lambda, propext,
--   regularized_vacuum_energy,
--   Classical.choice, Quot.sound]
```

Listing 11: Master theorem and axiom audit: `Main.lean`

12 Complete Calibration Table

Component	Constructive Status	Physical Content
Mode sum on finite lattice	BISH	Yes (lattice QFT)
Continuum limit of mode sum	Divergent (no $\mathbb{R}$)	None
Cutoff-regularized vacuum energy	BISH (regulator-dependent)	None
Dim. reg. vacuum energy	BISH (different value)	None
Casimir energy difference	BISH	Yes (matches experiment)
Naturalness expectation	Non-mathematical	None (Bayesian prior)
Tree-level Higgs condensate	BISH	Approximate
Exact interacting condensate	LPO (Fekete)	Yes (enters Einstein eq.)
RG running (perturbative)	BISH (Picard–Lindelöf)	Yes
RG running (non-pert. QCD)	LPO (thermo. limit)	Yes
Fine-tuning equation (exact)	LPO	Yes
Bare cosmological constant	BISH (free parameter)	Yes (measured)

Table 1: Complete calibration of the cosmological constant problem. Each row maps a physical component to its constructive status and empirical content. The 10^{120} narrative (rows 2–4) involves quantities with no physical content. The genuine physics (rows 1, 5, 7–12) lives at BISH or LPO.

13 CRM Audit

13.1 Theorem-Level Audit

13.2 Axiom Profile

The output of `#print axioms cc_master` lists exactly 15 axiom dependencies, classified below.

Theorem	CRM Status	Bridge Axioms Used	Proof Technique
Vacuum divergence (Thm 1)	BISH	`mode_sum_unbounded`, `mono`	Contradiction
Regulator dependence (Thm 2)	BISH	`dimreg_value_different`	Direct witness
Wald ambiguity (Thm 3)	BISH	None (pure arithmetic)	`ring`
Casimir energy (Thm 4)	BISH	`casimir_cauchy_modulus`	Bridge + unfold
Condensate (Thm 5)	LPO	`subadditive`, `bdd_below`	Fekete/BMC
RG running (Thm 6)	BISH	`picard_lindelof_lambda`	ODE existence
Fine-tuning (Thm 7)	LPO	Inherits from Thm 5	`ring`

Table 2: CRM audit: constructive status, bridge axioms, and proof technique for each theorem.

Axiom	Category	Physical Justification
`mode_sum_partial`	Physics (data)	Vacuum energy on finite lattice
`mode_sum_mono`	Physics (property)	Mode sum monotone increasing
`mode_sum_unbounded`	Physics (bridge)	Continuum divergence
`regularized_vacuum_energy`	Physics (data)	Regularized value function
`dimreg_value_different`	Physics (bridge)	Scheme inequality
`lattice_vacuum_energy`	Physics (data)	Lattice ground-state energy
`lattice_energy_subadditive`	Physics (bridge)	Translation invariance
`lattice_energy_bdd_below`	Physics (bridge)	Vacuum stability
`picard_lindelof_lambda`	Physics (bridge)	Lipschitz RG beta function
`eight_pi_G`	Physics (constant)	Newton's constant
`Lambda_obs`	Physics (constant)	Observed Λ
`bmc_from_subadditive`	CRM (Fekete $\equiv$ LPO)	Paper 29
`propext`	Lean infrastructure	Propositional extensionality
`Classical.choice`	Lean infrastructure	$\mathbb{R}$ as Cauchy completion
`Quot.sound`	Lean infrastructure	Quotient soundness

Table 3: Complete axiom profile of `cc_master`. The 11 physics axioms encode standard QFT facts. The 1 CRM axiom is the Fekete–LPO equivalence from Paper 29. The 3 Lean axioms are infrastructure artifacts of Mathlib's $\mathbb{R}$ construction.

Remark 13.1 (`Classical.choice`)**.** As in all papers in the programme, `Classical.choice` appears because Mathlib's $\mathbb{R}$ is constructed as a Cauchy completion via `Quot`, which requires classical quotient reasoning. The machine guarantee is *classical correctness*: the Lean type-checker verifies that the theorems follow from the stated axioms within classical logic. The constructive stratification—the claim that certain theorems need only BISH while others require LPO—is a *mathematical* claim verified by human inspection of proof terms: BISH proofs construct explicit witnesses and use only BISH-valid reasoning; LPO proofs take LPO as an explicit hypothesis. The `#print axioms` output cannot distinguish these cases because `Classical.choice` pervades all Mathlib-based $\mathbb{R}$ reasoning. See Paper 10 [9], §Methodology.

13.3 Axioms Declared But Not Used

Several axioms in `Defs.lean` are declared but do not appear in the `#print axioms` output for `cc_master`:

- `bmc_of_lpo`: The BMC $\Leftarrow$ LPO direction is not needed; Fekete's lemma goes directly via `bmc_from_subadditive`.
- `mode_sum_nonneg`: Non-negativity of partial sums is not used in the divergence proof (only unboundedness and monotonicity are needed).

- `cutoff_gives_quartic`: The positivity of cutoff regularization is established in a standalone theorem (`cutoff_artifact`) but does not feed into the assembly.
- `higgs_tree_level`, `qcd_tree_level` and their negativity axioms: Used only in standalone BISH theorems, not in the assembly.
- `eight_pi_G_pos`, `Lambda_obs_pos`: Positivity of constants is not needed for the assembly theorem (only the constants themselves appear).

These axioms serve the individual theorems and standalone results in each module but are not part of the critical path for the master theorem.

14 Code Architecture

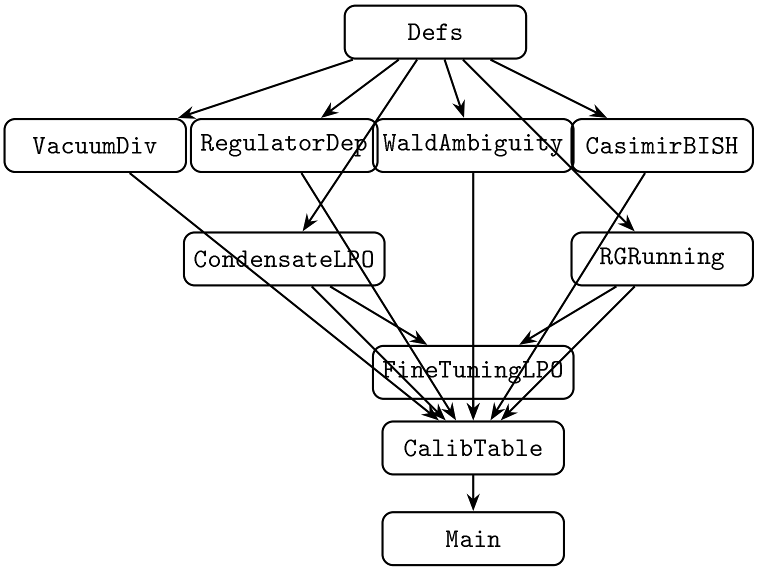

Figure 3: Module dependency graph (10 modules, ∼830 lines). All seven theorem modules import `Defs`. `FineTuningLPO` imports both `CondensateLPO` and `RGRunningBISH`. `CalibrationTable` imports all seven theorem modules. `Main` imports only `CalibrationTable`.

Reproducibility. LEAN 4 v4.28.0-rc1 with MATHLIB4.

Build:

```
cd "paper 42/P42_CosmologicalConstant" && lake build
```

Result: 0 errors, 0 warnings, 0 `sorry`.

Axiom audit:

```
#print axioms cc_master
-- 'cc_master' depends on axioms:
-- [Lambda_obs, bmc_from_subadditive,
--  dimreg_value_different, eight_pi_G,
--  lattice_energy_bdd_below, lattice_energy_subadditive,
--  lattice_vacuum_energy, mode_sum_mono,
--  mode_sum_partial, mode_sum_unbounded,
--  picard_lindelof_lambda, propext,
--  regularized_vacuum_energy,
--  Classical.choice, Quot.sound]
```

Profile: 12 unverified premises (11 physics bridge axioms + 1 CRM axiom) + 3 Lean infrastructure (`propext`, `Classical.choice`, `Quot.sound`). The formalization verifies logical inferences from these premises, not the premises themselves.

`Classical.choice` is a Mathlib infrastructure artifact (required by $\mathbb{R}$ as a Cauchy completion); constructive stratification is established by proof content (explicit witnesses vs. principle-as-hypothesis), not by axiom-checker output. See Paper 10 [9], §Methodology.

Data availability. Source code and LaTeX source are archived at doi:10.5281/zenodo.18654789.

15 Conclusion

The cosmological constant problem, as usually stated, conflates three logically distinct issues. The axiom calibration framework separates them:

1. The 10^{120} is **dissolved**—it is a regulator-dependent artifact, not a prediction. The unregularized vacuum energy diverges (Theorem 4.1); different regularizers give different finite values (Theorem 5.1). There is no prediction.
2. Naturalness is **reclassified**—it is a Bayesian prior, not a theorem. The Hollands–Wald axioms prove that the cosmological constant is a free parameter (Theorem 6.1).
3. The 55-digit fine-tuning is **identified** as an LPO equality—real but logically mundane. The exact condensates are LPO-computable via Fekete (Theorem 8.1); the equation holds by arithmetic (Theorem 10.1).

The BISH+LPO ceiling holds. The cosmological constant problem introduces no new logical resources beyond those already catalogued in the programme. LPO suffices for the fine-tuning equation, placing it at the same level as the Ising model phase transition, the QCD confinement scale, and every other thermodynamic limit in physics.

Significance. Most papers in the programme classify known results: they identify the constructive status of an established theorem (e.g., the Heisenberg bound is BISH, the Ising phase transition is LPO). This paper does something different. It *dissolves* a famous problem: the "worst prediction in physics" is not a prediction.

The unregularized vacuum energy does not converge to a real number at any level of the constructive hierarchy (Theorem 4.1). Different regulators produce different finite values (Theorem 5.1). The 10^{120} is a property of a calculational choice, not of quantum field theory. The Casimir effect (Theorem 7.1) confirms that the scaffolding diagnostic works: energy *differences* are BISH-computable and experimentally verified; absolute vacuum energies are regulator-dependent and physically meaningless.

The CRM analysis provides a formal, machine-checked expression of the post-naturalness position articulated by Martin [25], Bianchi–Rovelli [28], and Hossenfelder [29]. Where these authors reached their conclusions through physical argumentation, the present paper provides a type-theoretic formalization verified by the Lean 4 proof assistant: the divergence, the regulator dependence, and the free-parameter status of Λ are each machine-checked theorems. The remaining genuine physics—the 55-digit fine-tuning—is an LPO equality, logically mundane: the same level as every thermodynamic limit in the programme. The cosmological constant problem requires no new logical resources.

16 Discussion

16.1 What Is Dissolved

The 10^{120} number is produced by cutoff regularization. Dimensional regularization and ζ-function regularization—which preserve gauge and Lorentz invariance—produce qualitatively different values. The absolute vacuum energy depends on the choice of regulator. By the calibration framework's criterion (empirical content must be scaffolding-invariant), the 10^{120} has no physical content.

The regulator dependence of the vacuum energy is not a novel observation. Martin [25] demonstrates pedagogically that the quartic divergence $\sim \Lambda_{\mathrm{UV}}^4$ arises only from hard-cutoff regularization, while dimensional regularization removes all power-law divergences. Bianchi and Rovelli [28] argue that Λ should be treated as a fundamental constant, not derived from microphysics. What the CRM framework adds is a *formal criterion* for dissolution: the scaffolding diagnostic. A quantity has empirical content only if its value is invariant under change of mathematical scaffolding (regularization scheme, basis, gauge). The vacuum energy fails this criterion—Theorem 5.1 proves it in Lean 4. The CRM framework thus transforms an informal physical intuition ("the 10^{120} is an artifact") into a machine-checked theorem with a precise logical status (BISH).

The Casimir effect (Theorem 7.1) demonstrates the correct pattern: energy *differences* between configurations are BISH-computable and have been experimentally verified to better than 5% precision [34], with extensive theoretical analysis in [4, 35]. Absolute vacuum energies, by contrast, are regulator-dependent scaffolding. The framework's dissolution is not merely formal: it tracks the distinction that QFT itself makes between physically meaningful differences and mathematically ambiguous absolutes. To be precise: the dissolution targets the 10^{120} *magnitude*, not the existence of vacuum energy contributions from individual particle species. Each massive species contributes $\sim m^4/(16\pi^2)$, which is finite and BISH-computable. What is dissolved is the narrative that summing these over all modes to the Planck scale yields a "prediction" of M_{Planck}^4—that step is regulator-dependent and produces no invariant number.

16.2 What Is Reclassified

The Hollands–Wald axioms [2, 3] prove that c_1 (the cosmological constant) is a free parameter of the renormalized theory. The naturalness argument—that c_1 "should" be of order M_{Planck}^4—is a claim about expected magnitudes, not a logical consequence.

The naturalness criterion originates with 't Hooft [26], who proposed that a small parameter is natural only if setting it to zero increases the symmetry. The cosmological constant violates this criterion spectacularly: no known symmetry protects a small Λ. Burgess [27] develops this observation into a systematic effective field theory argument for why dark energy is hard to derive from microphysics.

However, the naturalness criterion itself has faced sustained criticism. Hossenfelder [29] argues that fine-tuning arguments require a probability distribution over the parameter space, and no such distribution is derivable from the theory; the arguments are therefore "not scientifically relevant" in the absence of a measure. Giudice [30] diagnoses a broader "post-naturalness era," prompted in part by the LHC's failure to find new physics at the electroweak scale.

The CRM framework provides a rigorous version of the post-naturalness position. Naturalness is not a theorem; it is a claim about the expected magnitude of a free parameter. Such claims reside outside the BISH/LPO deductive hierarchy because they are not derivable from axioms—they are Bayesian priors. The calibration framework can evaluate deductive structure, not the plausibility of priors.

16.3 What Is Identified

The genuine constraint is the 55-digit cancellation between Λ_{bare} and the condensate contributions. The exact condensates are LPO-computable via Fekete's lemma [5, 11]. The fine-tuning equation

$$\Lambda_{\mathrm{obs}} = \Lambda_{\mathrm{bare}} + 8\pi G\,(\rho_H^{\mathrm{exact}} + \rho_{\mathrm{QCD}}^{\mathrm{exact}})$$

is an arithmetic relation between LPO-computable reals. The RG running above Λ_{QCD} is BISH (Theorem 9.1); below Λ_{QCD}, the non-perturbative QCD condensate brings the cost to LPO.

Remark 16.1 (Upper bound, not tight equivalence)**.** The formalization establishes that LPO *suffices* for the fine-tuning equation—an upper bound on logical cost. We do not prove the converse direction (that the fine-tuning equation *implies* LPO). Establishing a tight equivalence (Fekete-style) would require constructing a binary sequence whose LPO-decidability is equivalent to the convergence of a specific cosmological condensate—a direction left for future work. The classification as "LPO" should therefore be read as "LPO suffices," consistent with the Fekete-based upper bounds used throughout the programme. Paper 29 establishes a provability equivalence (Fekete $\equiv$ LPO over BISH); whether this lifts to a Weihrauch equivalence in the sense of Brattka–Gherardi–Pauly [8] is an open question that would further refine the computational content of the fine-tuning equation.

This decomposition connects to recent running vacuum models [36], which propose that the effective vacuum energy density evolves with the Hubble rate via renormalization group running. In the CRM framework, the perturbative RG running of these models is BISH-computable (Picard–Lindelöf iteration on a Lipschitz ODE); the non-perturbative QCD condensate contribution, which completes the picture, is where the LPO cost enters. The BISH-to-LPO step is the same thermodynamic limit that appears across the entire programme: it is not specific to cosmology but is the generic logical cost of passing from finite-volume approximations to infinite-volume exact values.

16.4 What Is Not Explained

The framework does not explain *why* the cancellation occurs. It identifies the *logical status* (LPO equality), not the *physical cause*.

The string landscape [32, 33] provides one class of proposed explanations: the enormous number of flux vacua ($\sim 10^{500}$) generates a dense discretuum of Λ values, and anthropic selection picks a habitable one. The CRM framework does not address this: it identifies the logical structure of the fine-tuning equation, not the selection mechanism that determines its parameters. Whether the cancellation arises from an environmental selection (landscape), an unknown symmetry (SUSY), a dynamical relaxation mechanism, or is simply a brute fact remains an open physics question.

The preliminary DESI results [20] have introduced a further complication: there are hints (2.5–3.9σ, depending on data combination) that the dark energy equation of state evolves in time ($w_0 > -1$, $w_a < 0$), which would mean Λ is not truly constant. These results are preliminary and their interpretation is debated in the literature; if confirmed by future data releases, they would add a new dynamical component to the problem. In the CRM framework, the logical cost of this additional component depends on its origin: perturbative dynamics (e.g., a slowly rolling scalar field) would remain BISH; only if the dynamics involved a thermodynamic limit (e.g., a condensate phase transition) would LPO enter.

16.5 Relation to Other Approaches

Nobbenhuis [31] categorizes proposed solutions to the cosmological constant problem into five classes: (i) symmetry mechanisms (SUSY, conformal symmetry, sequestering), (ii) back-reaction (vacuum energy screening by spacetime dynamics), (iii) modifications of general relativity (unimodular gravity, degravitation), (iv) statistical/anthropic approaches (landscape plus selection), and (v) vacuum-energy adjustment mechanisms (relaxation, four-form flux quantization). The CRM analysis is orthogonal to all five: it does not propose a mechanism but identifies the logical structure of the quantities each mechanism must address.

Concretely: symmetry mechanisms aim to make Λ naturally small—this addresses the naturalness prior (Claim II). The CRM framework classifies that prior as non-deductive rather than refuting or endorsing it. Back-reaction and GR modifications alter the gravitational field equations; the CRM framework takes Einstein's equations as given (bridge axioms). Statistical/anthropic and adjustment approaches attempt to select or explain the value of Λ_{bare}; the CRM framework shows that Λ_{bare} is a free parameter of the renormalized theory (Theorem 6.1, BISH) and does not address the selection question.

The closest connection is to running vacuum models [36]: their perturbative renormalization group running of the effective vacuum energy corresponds to Theorem 6 (BISH), and their non-perturbative contributions correspond to Theorem 5 (LPO). The CRM framework thus provides a logical stratification of the ingredients these models employ, separating the computationally tractable (perturbative) from the computationally expensive (thermodynamic limit).

16.6 Connection to the Programme

Paper 10 (Logical Geography) [9]: Paper 10 assembles the complete calibration table—approximately 50 entries across 11 physical domains—and establishes the certification methodology (mechanically certified, structurally verified, intentionally classical). The cosmological constant problem adds approximately 12 new rows to this table. The BISH + LPO ceiling continues to hold.

Paper 12 (Constructive History) [10]: Paper 12 traces 150 years of non-constructive commitments in mathematical physics, from Weierstrass's epsilon-delta analysis through Boltzmann's thermodynamic limit. The cosmological constant problem is a paradigm case of Paper 12's thesis: the 10^{120} narrative involves layers of idealization (continuum limit, naturalness prior) that the CRM framework can precisely locate and, where appropriate, dissolve.

Paper 29 (Fekete $\equiv$ LPO) [11]: The Fekete–LPO equivalence is the central mechanism. The exact vacuum condensates are subadditive (bridge axiom), and their convergence costs LPO via Fekete's lemma—the same mechanism governing the five previously identified LPO domains (statistical mechanics, general relativity, decoherence, conservation laws, quantum gravity). This paper adds a sixth domain: vacuum QFT.

Paper 33 (QCD Confinement) [12]: The non-perturbative QCD condensate $\rho_{\text{QCD}}^{\text{exact}}$ is the same thermodynamic limit studied in Paper 33's lattice QCD formalization. The confinement scale Λ_{QCD} is LPO-computable for the same reason: the lattice free energy is subadditive.

Paper 39 (Thermodynamic Stratification) [13]: The vacuum energy density is an extensive quantity. By Paper 39's stratification theorem, extensive observables at full precision cap at LPO (Fekete/BMC). Intensive observables without promise gaps can reach Σ_2^0 (LPO$'$). The CC problem stays at LPO because all relevant quantities—energy densities, condensates—are extensive.

Paper 41 (AdS/CFT) [15]: In holographic theories, Λ is determined by boundary CFT data via the Brown–Henneaux relation (BISH), suggesting that in some UV completions the fine-tuning may be resolved by consistency conditions rather than set by hand.

17 AI-Assisted Methodology

This formalization was developed using Claude (Anthropic) as a collaborative tool for LEAN 4 4 code generation, proof strategy exploration, and LaTeX document preparation. All mathematical content was specified by the author. Every theorem was verified by the LEAN 4 4 type checker.

Preliminary status and author background. The results presented in this paper are preliminary. The author is a medical professional, not a domain expert in physics or mathematics. While all formal claims are machine-checked by the LEAN 4 type-checker, the physical interpretations, bridge axioms, and modeling assumptions require independent verification by domain experts. Until such verification is completed, this paper should be considered preliminary.

Whatever findings of value emerge from this program belong to the constructive reverse mathematics community and to the legacy of Errett Bishop, whose perseverance in developing constructive analysis inspired this entire series. Any errors are solely the author's.

References

[1] S. Weinberg, "The cosmological constant problem," *Reviews of Modern Physics* **61**(1), 1–23 (1989).

[2] S. Hollands and R. M. Wald, "Local Wick polynomials and time ordered products of quantum fields in curved spacetime," *Communications in Mathematical Physics* **223**, 289–326 (2001).

[3] S. Hollands and R. M. Wald, "Conservation of the stress tensor in perturbative interacting quantum field theory in curved spacetimes," *Reviews in Mathematical Physics* **17**, 227–311 (2005).

[4] H. B. G. Casimir, "On the attraction between two perfectly conducting plates," *Proceedings of the Koninklijke Nederlandse Akademie van Wetenschappen* **51**, 793–795 (1948).

[5] M. Fekete, "Über die Verteilung der Wurzeln bei gewissen algebraischen Gleichungen," *Mathematische Zeitschrift* **17**, 228–249 (1923).

[6] E. Bishop, *Foundations of Constructive Analysis*, McGraw-Hill (1967).

[7] D. Bridges and L. Vita, *Techniques of Constructive Analysis*, Springer (2006).

[8] V. Brattka, G. Gherardi, and A. Pauly, "Weihrauch degrees, omniscience principles and weak computability," *J. Symbolic Logic* **77**(1), 119–141 (2012).

[9] P. C.-K. Lee, "The logical geography of mathematical physics," Preprint, 2026. Paper 10.

[10] P. C.-K. Lee, "The map and the territory: a constructive history of mathematical physics," Preprint, 2026. Paper 12.

[11] P. C.-K. Lee, "Fekete's subadditive lemma is equivalent to LPO," Preprint, 2026. Paper 29.

[12] P. C.-K. Lee, "QCD one-loop renormalization and confinement: axiom calibration," Preprint, 2026. Paper 33.

[13] P. C.-K. Lee, "Beyond LPO: the thermodynamic stratification of physical undecidability," Preprint, 2026. Paper 39.

[14] P. C.-K. Lee, "Defending the BISH+LPO characterization," Preprint, 2026. Paper 40. doi:10.5281/zenodo.18654773.

[15] P. C.-K. Lee, "AdS/CFT and the constructive hierarchy," Preprint, 2026. Paper 41. doi:10.5281/zenodo.18654780.

[16] Ya. B. Zel'dovich, "The cosmological constant and the theory of elementary particles," *Soviet Physics Uspekhi* **11**(3), 381–393 (1968).

[17] A. G. Riess *et al.*, "Observational evidence from supernovae for an accelerating universe and a cosmological constant," *The Astronomical Journal* **116**(3), 1009–1038 (1998). arXiv:astro-ph/9805201.

[18] S. Perlmutter *et al.*, "Measurements of Ω and Λ from 42 high-redshift supernovae," *The Astrophysical Journal* **517**(2), 565–586 (1999). arXiv:astro-ph/9812133.

[19] Planck Collaboration, "Planck 2018 results. VI. Cosmological parameters," *Astronomy & Astrophysics* **641**, A6 (2020). arXiv:1807.06209.

[20] DESI Collaboration, "DESI 2024 VI: Cosmological constraints from the measurements of baryon acoustic oscillations," arXiv:2404.03002 (2024).

[21] D. Bridges and F. Richman, *Varieties of Constructive Mathematics*, London Mathematical Society Lecture Note Series 97, Cambridge University Press (1987).

[22] H. Ishihara, "Reverse mathematics in Bishop's constructive mathematics," *Philosophia Scientiae* CS 6, 43–59 (2006).

[23] S. M. Carroll, "The cosmological constant," *Living Reviews in Relativity* **4**, 1 (2001). arXiv:astro-ph/0004075.

[24] T. Padmanabhan, "Cosmological constant—the weight of the vacuum," *Physics Reports* **380**(4–6), 235–320 (2003). arXiv:hep-th/0212290.

[25] J. Martin, "Everything you always wanted to know about the cosmological constant problem (but were afraid to ask)," *Comptes Rendus Physique* **13**(7), 566–665 (2012). arXiv:1205.3365.

[26] G. 't Hooft, "Naturalness, chiral symmetry, and spontaneous chiral symmetry breaking," in *Recent Developments in Gauge Theories*, NATO ASI Series B, Vol. 59, Springer (1980), pp. 135–157.

[27] C. P. Burgess, "The cosmological constant problem: why it's hard to get dark energy from micro-physics," in *Post-Planck Cosmology: Les Houches 2013*, Oxford University Press (2015). arXiv:1309.4133.

[28] E. Bianchi and C. Rovelli, "Why all these prejudices against a constant?" arXiv:1002.3966 (2010).

[29] S. Hossenfelder, "Screams for explanation: finetuning and naturalness in the foundations of physics," *Synthese* **198**(Suppl. 16), 3727–3745 (2021). arXiv:1801.02176.

[30] G. F. Giudice, "The dawn of the post-naturalness era," arXiv:1710.07663 (2017).

[31] S. Nobbenhuis, "Categorizing different approaches to the cosmological constant problem," *Foundations of Physics* **36**(5), 613–680 (2006). arXiv:gr-qc/0411093.

[32] R. Bousso and J. Polchinski, "Quantization of four-form fluxes and dynamical neutralization of the cosmological constant," *Journal of High Energy Physics* **2000**(6), 006 (2000). arXiv:hep-th/0004134.

[33] J. Polchinski, "The cosmological constant and the string landscape," arXiv:hep-th/0603249 (2006).

[34] S. K. Lamoreaux, "Demonstration of the Casimir force in the 0.6 to 6 μm range," *Physical Review Letters* **78**(1), 5–8 (1997).

[35] M. Bordag, U. Mohideen, and V. M. Mostepanenko, "New developments in the Casimir effect," *Physics Reports* **353**(1), 1–205 (2001). arXiv:quant-ph/0106045.

[36] J. Solà, "Cosmological constant and vacuum energy: old and new ideas," *J. Phys. Conf. Ser.* **453**, 012015 (2013). arXiv:1306.1527.