

De Rham Decidability and DPT Completeness

The p -adic Precision Bound

Paper 59/60, Constructive Reverse Mathematics Series

Paul Chun-Kit Lee
New York University, Brooklyn, NY*

February 2026

Abstract

We formalize the p -adic precision bound $N_M = v_p(\det(1 - \varphi))$ for elliptic curves with good reduction. The key identity

$$N_M = v_p(1 - a_p + p) = v_p(\#E(\mathbb{F}_p))$$

shows that the precision lost in p -adic verification equals the p -adic valuation of the point count on the reduction. The Hasse bound $a_p^2 \leq 4p$ implies $\#E(\mathbb{F}_p) \geq 1$, so N_M is always well-defined.

We verify N_M for 24 entries across 4 elliptic curves ($X_0(11)$, $X_0(14)$, $X_0(15)$, 37a) at primes of good reduction. Four entries are anomalous ($N_M \geq 1$); the remaining 20 are generic ($N_M = 0$):

Curve	p	$\det(1 - \varphi)$	N_M	Note
$X_0(11)$	5	5	1	$5 \mid 5$
$X_0(14)$	3	6	1	$3 \mid 6, 9 \nmid 6$
$X_0(15)$	2	4	2	Max N_M in table
$X_0(15)$	7	7	1	$7 \mid 7, 49 \nmid 7$

“Axiom 5” (de Rham decidability) is a theorem for geometric representations: de Rham (Faltings/Tsuji) \Rightarrow potentially semistable (Berger) \Rightarrow weakly admissible (Colmez–Fontaine) $\Rightarrow N_M$ computable in BISH.

CRM classification: BISH. All arithmetic is exact over \mathbb{Z} ; no omniscience principles invoked.

Lean 4 formalization: 6 modules (~ 762 lines), zero errors, zero warnings, zero sorry. Zero custom axioms; all precision bounds verified by `norm_num` and `simp`.

1 Introduction

1.1 Main results

This paper quantifies the p -adic precision cost of de Rham verification for elliptic curves with good reduction. The entire computation is pure integer arithmetic: no p -adic analysis or period rings appear in the Lean formalization.

- (A) **Hasse implies positivity** (Theorem 3.1). For $p \geq 2$ with $a_p^2 \leq 4p$: $1 - a_p + p > 0$. Equivalently, $\#E(\mathbb{F}_p) \geq 1$, so N_M is always well-defined.

*Lean 4 source code and reproducibility materials: <https://doi.org/10.5281/zenodo.18735931>

- (B) **Ordinary, non-anomalous** (Theorem 4.2). If $p \nmid (1 - a_p)$, then $N_M = 0$: no precision loss.
- (C) **Supersingular, no precision loss** (Theorem 4.1). When $a_p = 0$: $\det(1 - \varphi) = 1 + p$ and $N_M = 0$.
- (D) **Verification table** (§5). 24 entries across 4 elliptic curves ($X_0(11)$, $X_0(14)$, $X_0(15)$, 37a). Four anomalous ($N_M \geq 1$), 20 generic ($N_M = 0$). Maximum $N_M = 2$ at $X_0(15)/p=2$.
- (E) **Semistable extension** (Theorem 6.1). For split multiplicative reduction (Tate curves), $N_M^{\text{st}} = v_p(q_E)$ is computable from the j -invariant.
- (F) **DPT completeness** (Theorem 9.1). Axioms 1–3 plus de Rham decidability suffice for the decidability of numerical equivalence on $\text{CH}^*(X)$. No mixed motive axiom is needed: Ext^1 is invisible to numerical equivalence. The pure motive program (Papers 50–59) is complete.
- (G) **Analytic rank stratification** (Theorem 9.2). The logical complexity of computing $\text{Ext}^1(\mathbb{Q}(0), M)$ is stratified by the analytic rank $r = \text{ord}_{s=s_0} L(M, s)$: $r = 0$ and $r = 1$ are BISH; $r \geq 2$ requires MP.

1.2 Constructive reverse mathematics primer

Bishop-style constructive mathematics (BISH) [2] works within intuitionistic logic: no excluded middle, no axiom of choice. Classical theorems are recovered by adding *omniscience principles*:

$$\text{BISH} \subset \text{BISH} + \text{MP} \subset \text{BISH} + \text{LLPO} \subset \text{BISH} + \text{LPO} \subset \text{CLASS}.$$

Constructive reverse mathematics (CRM) classifies theorems by the *weakest* principle required. This paper operates entirely in BISH: all arithmetic is exact over \mathbb{Z} , all witnesses are explicit, and no omniscience principle is needed.

For the BISH/LPO/CLASS hierarchy and its role in physics, see the series overview (Paper 45 [8]).

1.3 The DPT framework at finite primes

Paper 50 [9] introduced the *Decidable Polarized Tannakian* (DPT) category as a constructive proxy for Grothendieck’s conjectural category of motives. The three axioms—**Axiom 1** (decidable morphisms, Standard Conjecture D), **Axiom 2** (algebraic spectrum), and **Axiom 3** (Archimedean polarization)—provide BISH-decidability at the infinite place via positive-definiteness ($u(\mathbb{R}) = 1$).

At finite primes, $u(\mathbb{Q}_p) = 4$ blocks positive-definiteness (Papers 45, 51 [8, 10]). The replacement strategy is *de Rham descent*:

- (1) V is de Rham (automatic for geometric representations; Faltings [4]/Tsuji [7]).
- (2) V is potentially semistable (Berger’s [1] equivalence: de Rham \Leftrightarrow pot. semistable).
- (3) $D_{\text{cris}}(V)$ is weakly admissible (Colmez–Fontaine [3]).
- (4) Weak admissibility provides the precision bound N_M .

Therefore “Axiom 5” is a *theorem* for motives, not an independent axiom. This paper quantifies the precision cost: how many extra p -adic digits are needed for verification? The answer is $N_M = v_p(\#E(\mathbb{F}_p))$.

1.4 From DPT to the precision formula: the trajectory

The formula $N_M = v_p(\#E(\mathbb{F}_p))$ was not computed in isolation. It emerged from a systematic exploration of the DPT framework’s boundary, following two complementary threads:

Thread 1: Axiom 1 boundary (codimension ≥ 2). Papers 54–55 [11, 12] identified the codimension principle: Axiom 1 holds at codimension 1 (Lefschetz) but fails at codimension ≥ 2 (exotic Hodge classes). Paper 56 [13] followed this prediction to exotic Weil classes on CM abelian fourfolds, discovering the formula $\deg(w_0 \cdot w_0) = \sqrt{\text{disc}(F)} = f$. Paper 57 [14] completed the enumeration across all nine Heegner fields and proved the cyclic barrier (the formula cannot extend to non-cyclic cubics). Paper 58 [15] extended the formula to $h_K > 1$ via the Steinitz class.

Thread 2: Axiom 3 boundary (p -adic, finite primes). Paper 51 [10] identified the p -adic BSD exceptional zero as the p -adic avatar of the u -invariant obstruction ($u(\mathbb{Q}_p) = 4$). This paper (Paper 59) quantifies the precision cost at this boundary: $N_M = v_p(\#E(\mathbb{F}_p))$ bounds the extra digits needed.

Together, these threads show that *both boundaries of the DPT landscape have computable structure* (Figure 1). The Axiom 1 boundary yields a topological invariant (the conductor f); the Axiom 3 boundary yields an arithmetic invariant (the p -adic valuation of the point count).




Figure 1: The two principal boundaries of the DPT landscape. Thread 1 (Axiom 1 boundary, codimension ≥ 2) yields the conductor formula for exotic Weil classes (Papers 56–58). Thread 2 (Axiom 3 boundary, p -adic) yields the precision formula (this paper). Both boundaries have computable arithmetic structure.

1.5 Connection to Paper 51

Paper 51 [10] identified the p -adic BSD exceptional zero as the p -adic avatar of the u -invariant obstruction. Paper 59 quantifies this:

- $N_M = 0$: standard decidability, no anomaly.
- $N_M \geq 1$: precision loss equals the “exceptional” extra computation.
- The \mathcal{L} -invariant compensates for this precision loss.

The exceptional zero is not a failure of decidability—it is a quantified precision cost, bounded by N_M .

1.6 State of the art

The p -adic comparison theorems of Faltings [4] and Tsuji [7] establish that geometric Galois representations are de Rham. Berger [1] proved the equivalence de Rham \Leftrightarrow potentially semistable. Colmez–Fontaine [3] proved weakly admissible \Leftrightarrow admissible, completing the chain from de Rham to concrete filtered φ -module data. The MTT exceptional zero conjecture [5] relates the p -adic L -function to the \mathcal{L} -invariant at anomalous primes.

The numerical consequence of this chain—the precision bound $N_M = v_p(\det(1 - \varphi))$ —is elementary to compute for any given curve and prime. This paper records the computation for 24 entries across 4 curves, extracts the identity $N_M = v_p(\#E(\mathbb{F}_p))$, and uses this to answer a question internal to the CRM program: whether the DPT framework’s p -adic boundary requires a new axiom (§9).

1.7 Caveats

- (i) This paper does not formalize the chain de Rham \Rightarrow pot. semistable \Rightarrow weakly admissible. These deep theorems of p -adic Hodge theory are stated as mathematical background, not formalized in Lean.
- (ii) It does not compute Tate parameters for semistable reduction. The semistable extension (Theorem E) is a structural declaration.
- (iii) The a_p values are hardcoded from LMFDB/standard tables. The formalization verifies the *consequences* (precision bounds), not the point-counting algorithms that produce a_p .
- (iv) Extension to abelian varieties of dimension $g > 1$ is not attempted.
- (v) The computation of $N_M = v_p(1 - a_p + p)$ is elementary integer arithmetic; no number theorist would find the formula itself surprising. The Hasse bound is 1930s mathematics, the point count identity is undergraduate, and the p -adic Hodge theory chain (Faltings \rightarrow Berger \rightarrow Colmez–Fontaine) is cited, not proved. The contribution is not the arithmetic but the *framing*: closing Paper 54’s open question about whether the DPT framework requires a new axiom at its p -adic boundary.

2 Preliminaries

2.1 Crystalline representations of elliptic curves

Let E/\mathbb{Q} be an elliptic curve with good reduction at a prime p . The p -adic Galois representation $V = V_p(E)$ is crystalline at p , and the filtered φ -module $D_{\text{cris}}(V)$ is 2-dimensional over \mathbb{Q}_p . The Frobenius φ has characteristic polynomial

$$\det(T \cdot \text{Id} - \varphi) = T^2 - a_p T + p,$$

where $a_p = \text{Tr}(\varphi)$ is the trace of Frobenius [6].

2.2 The Hasse bound

The Hasse bound states $|a_p| \leq 2\sqrt{p}$. In squared form (avoiding irrationals):

$$a_p^2 \leq 4p. \tag{1}$$

This is the form used in the Lean formalization.

2.3 The precision bound

The key operator $(1 - \varphi) : D_{\text{cris}}(V) \rightarrow D_{\text{cris}}(V)$ has determinant

$$\det(1 - \varphi) = (1 - \alpha)(1 - \beta) = 1 - (\alpha + \beta) + \alpha\beta = 1 - a_p + p \quad (2)$$

where α, β are the eigenvalues of φ .

The p -adic precision bound is

$$N_M = v_p(\det(1 - \varphi)) = v_p(1 - a_p + p). \quad (3)$$

To verify $x = 0$ in $D_{\text{cris}}(V)$ to precision k , it suffices to compute modulo p^{k+N_M} .

2.4 The point count identity

Since $\#E(\mathbb{F}_p) = 1 - a_p + p$, we have the identity:

$$N_M = v_p(\#E(\mathbb{F}_p)). \quad (4)$$

The precision lost in p -adic verification equals the p -adic valuation of the number of \mathbb{F}_p -rational points on the reduction.

3 Hasse bound implies positivity

Theorem 3.1 (Theorem A: Hasse implies positivity). *For $p \geq 2$ and $a_p^2 \leq 4p$:*

$$1 - a_p + p > 0.$$

Equivalently, $\#E(\mathbb{F}_p) \geq 1$.

Proof. If $a_p \geq p + 1$, then $a_p^2 \geq (p + 1)^2 = (p - 1)^2 + 4p > 4p$ (since $(p - 1)^2 > 0$ for $p \geq 2$), contradicting the Hasse bound. Therefore $a_p \leq p$, and $1 - a_p + p \geq 1 > 0$.

In Lean, this is closed by `nlinarith` with the hint `sq_nonneg (a_p - (p : ℤ) - 1)`. \square

Corollary 3.2. $N_M = v_p(\#E(\mathbb{F}_p))$ is well-defined for all primes of good reduction.

4 Case analysis

4.1 Supersingular case

Theorem 4.1 (Theorem C: Supersingular, no precision loss). *When $a_p = 0$ (supersingular reduction for $p \geq 5$):*

$$\det(1 - \varphi) = 1 + p, \quad N_M = 0.$$

Proof. We must show $p \nmid (1 + p)$. If $p \mid (1 + p)$, then $p \mid 1$, so $p \leq 1$, contradicting $p \geq 2$. \square

4.2 Ordinary, non-anomalous case

Theorem 4.2 (Theorem B: Ordinary, non-anomalous). *If $p \nmid (1 - a_p)$, then $p \nmid (1 - a_p + p)$, so $N_M = 0$.*

Proof. If $p \mid (1 - a_p + p)$, then since $p \mid p$, we get $p \mid (1 - a_p)$, contradicting the hypothesis. \square

4.3 Anomalous case

When $p \mid (1 - a_p)$, equivalently $p \mid \#E(\mathbb{F}_p)$, the prime is *anomalous* and $N_M \geq 1$. These are the primes where the p -adic BSD exceptional zero phenomenon manifests (Paper 51 [10]).

Definition 4.3. A prime p of good reduction is *anomalous* for E if $(p : \mathbb{Z}) \mid \det(1 - \varphi)$, i.e., $N_M \geq 1$.

5 Verification table

All 24 entries are machine-verified in Lean. Each entry certifies:

- (i) $1 - a_p + p = \det$ (by `norm_num`).
- (ii) $N_M = v_p(\det)$ (by `simp`, checking divisibility).

5.1 $X_0(11)$ — Cremona label 11a1, conductor 11

$y^2 + y = x^3 - x^2 - 10x - 20$. Good reduction at all primes $\neq 11$.

p	a_p	$\det(1 - \varphi)$	N_M	Note
2	-2	5	0	
3	-1	5	0	
5	1	5	1	Anomalous: $5 \mid 5$
7	-2	10	0	
13	4	10	0	
17	-2	20	0	
19	0	20	0	Supersingular: $a_{19} = 0$
23	-1	25	0	

5.2 $X_0(14)$ — Cremona label 14a1, conductor 14

Bad reduction at 2 and 7.

p	a_p	$\det(1 - \varphi)$	N_M	Note
3	-2	6	1	Anomalous: $3 \mid 6, 9 \nmid 6$
5	-1	7	0	
11	-1	13	0	
13	4	10	0	
17	-2	20	0	

5.3 $X_0(15)$ — Cremona label 15a1, conductor 15

Bad reduction at 3 and 5.

p	a_p	$\det(1 - \varphi)$	N_M	Note
2	-1	4	2	Anomalous: $4 \mid 4, 8 \nmid 4$. Max N_M in table
7	1	7	1	Anomalous: $7 \mid 7, 49 \nmid 7$
11	-2	14	0	
13	4	10	0	
17	-2	20	0	

5.4 37a — Cremona label 37a1, conductor 37

$y^2 + y = x^3 - x$. Bad reduction only at 37.

p	a_p	$\det(1-\varphi)$	N_M	Note
2	-2	5	0	
3	-3	7	0	
5	-2	8	0	
7	-2	10	0	
11	0	12	0	Supersingular: $a_{11} = 0$
13	5	9	0	

5.5 Summary statistics

Statistic	Value
Total entries	24
Anomalous ($N_M \geq 1$)	4
Generic ($N_M = 0$)	20
Maximum N_M	2 ($X_0(15)$ at $p = 2$)
All $\det > 0$	✓ (Hasse bound)

All statistics are verified by `native_decide` in Lean.

6 Semistable extension

Theorem 6.1 (Theorem E: Semistable precision bound). *For E/\mathbb{Q} with split multiplicative reduction at p (Tate curve), $D_{\text{cris}}(V)$ does not exist. Instead, $D_{\text{st}}(V)$ is a 2-dimensional filtered (φ, N) -module with $N \neq 0$. The precision bound uses the Tate parameter q_E :*

$$N_M^{\text{st}} = v_p(q_E),$$

which is computable from the j -invariant: $v_p(q_E) = -v_p(j(E))$ when $v_p(j(E)) < 0$.

This extension shows that even at primes of bad (semistable) reduction, the precision bound is computable. The MTT exceptional zero (Paper 51 [10, 5]) has a quantified precision cost. This is formalized as a structural declaration (`SemistableData`) without proof, as computing Tate parameters requires p -adic analysis beyond integer arithmetic.

7 CRM audit

Classification: BISH.

1. **Hasse positivity.** The proof that $1 - a_p + p > 0$ uses `nlinarith` with algebraic hints. No omniscience needed; the Hasse bound converts the positivity check into bounded integer arithmetic.
2. **Case analysis.** The supersingular case uses direct divisibility ($p \mid (1 + p) \Rightarrow p \mid 1$, contradiction). The ordinary non-anomalous case uses $p \mid p$ and $p \nmid (1 - a_p)$ to conclude $p \nmid (1 - a_p + p)$. Both are purely algebraic.
3. **Verification table.** All 24 entries verified by `norm_num` (determinant computation) and `simp` (divisibility check). Summary statistics by `native_decide`.

4. **Precision bound computation.** Given E and p , the algorithm terminates in BISH: (i) look up a_p (finite), (ii) compute $1 - a_p + p$ (integer arithmetic), (iii) trial-divide by p to get N_M (finite).
5. **No omniscience.** No step invokes LPO, LLPO, MP, or WLPO.

8 Formal verification

The Lean 4 formalization builds with zero errors and zero warnings under `leanprover/lean4:v4.29.0-rc1` with Mathlib.

8.1 Module structure

#	Module	Lines	Sorry budget
1	Defs	74	0
2	PadicVal	54	0
3	VerificationTable	272	0
4	WeakAdmissibility	128	0
5	Interpretation	185	0
6	Main	49	0
Total		762	0 sorry, 0 custom axioms

8.2 Axiom inventory

Zero custom axioms. The precision bound $N_M = v_p(\#E(\mathbb{F}_p))$ is verified per-entry as a proof obligation in the `VerifiedBound` structure, not as a global axiom. Each entry proves its bound by `norm_num` and `simp`. All other theorems (Hasse positivity, supersingular case, ordinary case) are proved from the integer structure field hypotheses.

8.3 Code excerpts

Module 1: Core structures.

```
structure GoodReductionData where
  curve : EllipticCurveData
  p : N
  p_prime : Nat.Prime p
  a_p : Z
  hasse_bound : a_p ^ 2 ≤ 4 × (p : Z)

def det_one_minus_frob (d : GoodReductionData) : Z :=
  1 - d.a_p + d.p
```

Module 3: Verification table entry.

```
structure VerifiedBound where
  label : String; p : N; a_p : Z
  det_val : Z; N_M : N
  det_eq : (1 : Z) - a_p + p = det_val
  bound_correct :
    if N_M = 0 then ¬(p : Z) ∣ det_val
    else (p : Z) ^ N_M ∣ det_val
      ∧ ¬(p : Z) ^ (N_M + 1) ∣ det_val

def X0_11_at_5 : VerifiedBound where
  label := "11a1"; p := 5; a_p := 1
```

```

det_val := 5; N_M := 1
det_eq := by norm_num
bound_correct := by simp

```

Module 4: Hasse positivity.

```

theorem hasse_implies_positive (p : N) (a_p : Z)
  (hp : p ≥ 2) (hasse : a_p ^ 2 ≤ 4 × (p : Z)) :
  1 - a_p + (p : Z) > 0 := by
  nlinarith [sq_nonneg (a_p - (p : Z) - 1),
  sq_nonneg ((p : Z) - 1)]

```

Module 4: Supersingular case.

```

theorem supersingular_no_precision_loss (p : N) (hp : p ≥ 2) :
  ¬ (p : Z) ∣ (1 + (p : Z)) := by
  rintro ⟨k, hk⟩
  have h1 : (p : Z) ∣ 1 :=
    ⟨k - 1, by linarith [mul_sub (p : Z) k 1]⟩
  have h2 : (p : Z) ≤ 1 := Int.le_of_dvd one_pos h1
  linarith

```

Module 3: Summary statistics.

```

theorem table_length :
  verification_table.length = 24 := by native_decide

theorem anomalous_count :
  (verification_table.filter
    (fun v => decide (v.N_M > 0))).length = 4 := by
  native_decide

```

8.4 #print axioms output

All key theorems (`hasse_implies_positive`, `det_pos`, `supersingular_no_precision_loss`, `ordinary_non_anomalous`, `table_length`, `anomalous_count`) depend only on:

- `propext` (propositional extensionality—Lean kernel)
- `Classical.choice` (Mathlib infrastructure)
- `Quot.sound` (quotient soundness—Lean kernel)

No custom axioms appear.

8.5 Classical.choice audit

`Classical.choice` appears through Mathlib’s `Int.dvd_iff_emod_eq_zero` and related integer divisibility infrastructure. The proof *content* is entirely computational: all determinants are verified by `norm_num` on concrete integers, all divisibility checks by `simp`, and all summary statistics by `native_decide`. The constructive stratification is established by proof content (explicit witnesses, no omniscience hypotheses), not by axiom-checker output—following the methodology of Paper 10 (see Paper 45 [8]). The BISH classification is genuine at the proof-content level.

8.6 Reproducibility

- **Lean version:** `leanprover/lean4:v4.29.0-rc1` (pinned in `lean-toolchain`).
- **Mathlib:** resolved via `lakefile.lean` (commit pinned in `lake-manifest.json`).
- **Build:** `cd P59_DeRhamDecidability && lake build` produces zero errors, zero warnings, zero sorry.
- **Source:** <https://doi.org/10.5281/zenodo.18735931>

9 Discussion

9.1 Resolving Paper 54’s “Axiom 5” question

Paper 54 [11] identified two fracture points in the DPT framework where BISH-decidability breaks down:

- **Fracture Point 1** (Ext^1 boundary): Axiom 1 fails for mixed motives—extension classes in Ext^1 lie outside the Tannakian envelope, and the morphism decidability of Axiom 1 does not extend to them.
- **Fracture Point 2** (p -adic boundary): Axiom 3 fails at finite primes— $u(\mathbb{Q}_p) = 4$ blocks positive-definiteness, and the Archimedean polarization that secures decidability at the infinite place has no p -adic analogue.

Paper 54 posed the explicit open question: does p -adic decidability require a new “Axiom 5” to supplement the three axioms of Paper 50 [9]? See Figure 2 for a summary.

The answer: no, for geometric representations. The chain

de Rham (Faltings/Tsuji) \Rightarrow pot. semistable (Berger) \Rightarrow weakly admissible (Colmez–Fontaine)

reduces p -adic verification to filtered φ -module data, and the precision cost $N_M = v_p(\#E(\mathbb{F}_p))$ is computable in BISH by pure integer arithmetic. No new axiom is needed—the existing theorems of p -adic Hodge theory *are* the axiom.

Paper 54’s Tamagawa factor obstruction at finite primes—where local Tamagawa numbers introduce denominators that block constructive computation—dissolves for crystalline representations at primes of good reduction. The filtered φ -module $D_{\text{cris}}(V)$ provides the computable data directly, without passage through Tamagawa factors.

Partial resolution. Paper 59 heals Fracture Point 2 (p -adic boundary) for good reduction. Fracture Point 1 (Ext^1 boundary, mixed motives) remains open—this is the domain that Papers 56–58 [13, 14, 15] began to explore through exotic Weil classes on CM abelian fourfolds.

Two boundaries, two formulas. Both principal boundaries of the DPT landscape now have computable structure:

- The **Axiom 1 boundary** (codimension ≥ 2) yields a topological invariant: the conductor formula $\deg(w_0 \cdot w_0) = f$ across all nine Heegner fields (Paper 57 [14]), extended to $h_K > 1$ via the Steinitz class (Paper 58 [15]).
- The **Axiom 3 boundary** (p -adic, finite primes) yields an arithmetic invariant: the precision formula $N_M = v_p(\#E(\mathbb{F}_p))$ (this paper).

In both cases, the DPT framework’s prediction—that boundary objects carry computable arithmetic data—is confirmed by explicit computation.




Figure 2: Paper 54 identified two fracture points in the DPT framework and posed the question of whether an “Axiom 5” is needed. Paper 59 heals Fracture Point 2 for good reduction: no new axiom is required for geometric representations. Fracture Point 1 remains open.

9.2 The anomalous–generic dichotomy

The precision bound induces a clean dichotomy:

- **Generic** ($N_M = 0$, 20 of 24 entries): no precision loss, standard BISH-decidability.
- **Anomalous** ($N_M \geq 1$, 4 of 24 entries): quantified precision cost, bounded and computable. These are the primes where $p \mid \#E(\mathbb{F}_p)$, i.e., the p -adic BSD exceptional zero phenomenon manifests.

The dichotomy mirrors the cyclic barrier of Paper 57 [14]: both identify a clean structural boundary (cyclic vs. non-cyclic for Axiom 1; generic vs. anomalous for Axiom 3) where the transition is sharp.

9.3 Scope of contribution

The computation in this paper is elementary. The formula $N_M = v_p(1 - a_p + p)$ involves nothing beyond addition, subtraction, and trial division; the Hasse bound is classical, the point count identity standard, and the deep theorems of p -adic Hodge theory are cited, not formalized. Read in isolation, the paper verifies that certain well-known integers have certain well-known properties, and records this in Lean.

The contribution is not the arithmetic. It is the *completeness argument*: Paper 54 posed a specific question—does the DPT framework need a new axiom at its p -adic boundary?—and this paper provides a definite answer. The DPT axioms detect their own boundaries, and at those boundaries one finds computable invariants rather than voids. The verification table is the evidence; the theorem is that no “Axiom 5” is required.

More broadly, the CRM methodology—calibrating logical cost by the weakest principle needed—produces well-defined, machine-verifiable answers even at the framework’s boundaries. This is the sense in which the paper contributes: not by discovering new mathematics, but by demonstrating that the program’s internal questions have clean resolutions.

9.4 DPT completeness for numerical equivalence

With de Rham decidability established, the DPT framework for numerical equivalence is complete. Numerical equivalence on $\mathrm{CH}^k(X)$ is defined by the intersection pairing, which factors through the cycle class map to cohomology and is computed by traces of endomorphisms in the pure motive $h^*(X)$. The kernel of the cycle class map—homologically trivial cycles—is annihilated by the pairing and hence invisible to numerical equivalence. The extension groups $\mathrm{Ext}^1(M, N)$ in the mixed motive category govern this kernel (Abel–Jacobi images, Griffiths groups, Mordell–Weil groups), but since numerical equivalence projects away from the kernel, no Ext^1 computation is needed.

Theorem 9.1 (DPT completeness). *The decidability of numerical equivalence on $\mathrm{CH}^*(X)$ for smooth projective varieties over \mathbb{Q} requires only Axioms 1–3 of Paper 50 [9] plus de Rham decidability at finite primes (this paper). No mixed motive axiom is needed.*

This is a structural observation: the proof is the factorization above. The pure motive program (Papers 50–59) is now closed.

9.5 The mixed motive frontier: analytic rank stratification

The pure motive framework governs *numerical* equivalence. For *rational* equivalence—where the relevant invariant is $\mathrm{Ext}^1(\mathbb{Q}(0), M)$ —the logical complexity depends on the analytic rank $r = \mathrm{ord}_{s=s_0} L(M, s)$.

For a pure motive M , the group $\mathrm{Ext}^1(\mathbb{Q}(0), M)$ carries concrete arithmetic content:

- $M = h^1(E)$, E an elliptic curve: $\mathrm{Ext}^1 \cong E(\mathbb{Q}) \otimes \mathbb{Q}$ (Mordell–Weil group).
- $M = h^1(A)$, A an abelian variety: $\mathrm{Ext}^1 \cong A(\mathbb{Q}) \otimes \mathbb{Q}$.
- $M = h^2(X)(1)$, X a surface: Ext^1 relates to the Griffiths group.
- $M = \mathbb{Q}(n)$: Ext^1 relates to algebraic K -theory [20].

The CRM question: given the L -value $L^*(M, s_0)$, can one *constructively* extract generators of Ext^1 ?

Theorem 9.2 (Analytic rank stratification). *Let M be a motive over \mathbb{Q} possessing the Northcott property [19]. The logical complexity of constructively computing a basis for $\mathrm{Ext}^1(\mathbb{Q}(0), M)$ is stratified by r :*

r	Principle	Mechanism
0	BISH	$\mathrm{Ext}^1 = 0$; verify $L(M, s_0) \neq 0$ to finite precision
1	BISH	1-dim regulator; BK^* + Northcott bound the search
≥ 2	MP	Covolume $\not\propto$ basis vector bound (Minkowski)

*Conditional on Bloch–Kato [16], effective Silverman bounds [18], and effective lower bounds on $|L(M, s_0)|$.

Proof sketch. **Case $r = 0$.** Bloch–Kato predicts $\mathrm{Ext}^1 \otimes \mathbb{Q} = 0$. Verifying $L(M, s_0) \neq 0$ requires computing to precision below an effective lower bound on $|L(M, s_0)|$ —a bounded computation.

Case $r = 1$. The regulator $R(M) = \hat{h}(P)$ (the canonical height of the single generator) is determined by the Bloch–Kato formula. The Silverman height difference bound converts this to a naïve height bound; Northcott gives a finite, effectively enumerable search space.

Case $r \geq 2$. The regulator $R(M) = \det((P_i, P_j))$ is the Gram determinant of the Néron–Tate pairing. By Minkowski’s theorem on successive minima, $\lambda_1 \cdots \lambda_r \leq \gamma_r^{r/2} V$, bounding the product of successive minima but not the maximum. In dimension ≥ 2 , $\lambda_1 \rightarrow 0$ forces $\lambda_r \rightarrow \infty$ at fixed covolume. An enumeration by ascending height terminates (assuming finite III) but with no a priori bound—exactly MP.

Why MP and not LPO? The rank ≥ 2 computation is an unbounded search guaranteed to terminate (the generators exist and will eventually be found by ascending height enumeration). MP asserts exactly this: if a computation does not fail to terminate, it terminates. LPO would additionally provide a decision procedure for the order of vanishing itself—a strictly stronger requirement not needed here.

Promoting rank ≥ 2 to BISH would require Lang’s Height Lower Bound Conjecture (open). \square




Figure 3: The Minkowski obstruction. **Left:** In dimension 1, the lattice covolume is the generator length; knowing R determines P . **Right:** In dimension ≥ 2 , two lattices can share the same covolume V with arbitrarily different basis vector lengths. One real number cannot encode two independent heights.

Examples. $E = X_0(11)$ (rank 0): $L(E, 1) = 0.2538\dots \neq 0$; by Kolyvagin [21], $E(\mathbb{Q})$ is finite (in fact $\cong \mathbb{Z}/5\mathbb{Z}$); BISH. $E = 37a1$ (rank 1): $L'(E, 1) \neq 0$; by Gross–Zagier [17], the derivative determines a Heegner point height; generator $P = (0, 0)$ with $\hat{h}(P) = 0.0511\dots$; Northcott-bounded search; BISH. $E = 389a1$ (rank 2): regulator known but no bound on individual generator heights; MP.

Scope. The rank stratification applies CRM labels to standard material (Bloch–Kato, Silverman, Minkowski). The underlying arithmetic is not new. The contribution is the identification of analytic rank as the parameter governing the BISH/MP boundary for Ext^1 decidability. We note that identifying MP as the *exact* logical cost requires a reversal (Ext^1 computation implies MP over BISH), which we have not proved; the geometry-of-numbers obstruction makes such a reversal plausible but it remains open.

9.6 Open questions

Higher-dimensional representations. For abelian varieties of dimension $g > 1$, the crystalline Frobenius has $2g$ -dimensional characteristic polynomial. The precision bound generalizes to $N_M = v_p(\det(1 - \varphi))$, but the relationship to point counts is more complex.

Anomalous density. Among our 24 entries, 4 are anomalous (16.7%). The asymptotic density of anomalous primes for a given curve is predicted by Serre’s conjecture on Frobenius distributions.

\mathcal{L} -invariant connection. Paper 51 [10] identified the \mathcal{L} -invariant as the compensator for the exceptional zero. A precise formula connecting N_M to the p -adic BSD formula would close the loop between the precision bound and the \mathcal{L} -invariant.

Bad reduction beyond semistable. Theorem E addresses split multiplicative reduction via the Tate parameter q_E . For additive reduction, the filtered (φ, N, G_K) -module structure is more complex, and the precision bound requires potentially semistable (not just crystalline) data. A systematic computation of N_M for additive primes remains open.

Ext^1 fracture partial resolution. Paper 54’s Fracture Point 1 (Ext^1 boundary, mixed motives) is partially addressed by Papers 56–58 [13, 14, 15], which computed explicit invariants for exotic Weil classes on CM abelian fourfolds. However, a full constructive treatment of mixed motivic extensions—where Ext^1 classes resist Tannakian decidability—is unknown.

Lang’s Height Lower Bound Conjecture. An effective lower bound on the shortest vector λ_1 in the Mordell–Weil lattice would promote rank ≥ 2 from MP to BISH—bounding $\hat{h}(P)$ away from zero for non-torsion P .

Motives without Northcott. For K3 surfaces and higher K -theory, the relevant cycle groups lack a proven Northcott property. Without Northcott, even rank 1 computations are unbounded searches. The rank stratification (Theorem 9.2) applies only to motives with Northcott.

Higher Ext groups. The Ext^2 and higher groups in the mixed motive category are poorly understood. Their CRM classification is entirely open.

10 Conclusion

The p -adic precision bound $N_M = v_p(\#E(\mathbb{F}_p))$ is computable in BISH for all primes of good reduction. The Hasse bound guarantees $\#E(\mathbb{F}_p) \geq 1$, making N_M well-defined. The anomalous–generic dichotomy ($N_M = 0$ vs. $N_M \geq 1$) separates standard decidability from quantified precision cost.

Paper 54 [11] posed the question of whether the DPT framework requires a new “Axiom 5” for p -adic decidability. The answer, for geometric representations, is no: the existing theorems of p -adic Hodge theory reduce verification to elementary integer arithmetic, and no omniscience principle is needed. The resolution is partial—Fracture Point 2 (p -adic boundary) is healed for good reduction, while Fracture Point 1 (Ext^1 boundary, mixed motives) remains open—but it is definite.

The Lean 4 formalization (762 lines, zero sorry, zero custom axioms) verifies the arithmetic. The deeper claim is that the CRM methodology produces well-defined, machine-verifiable answers even at the framework’s boundaries. Together with Papers 56–58 (which found the conductor formula at the Axiom 1 boundary), this paper shows that both principal boundaries of the DPT landscape carry computable arithmetic invariants—not voids.

The pure motive program is complete: Axioms 1–3 plus de Rham decidability suffice for numerical equivalence, with no mixed motive axiom required (Theorem 9.1). The mixed motive frontier is stratified by analytic rank: $r = 0$ and $r = 1$ are BISH-decidable, while $r \geq 2$ requires Markov’s Principle—a structural obstruction imposed by the geometry of numbers, removable only by Lang’s Height Lower Bound Conjecture (Theorem 9.2).

Acknowledgments

We thank the Mathlib contributors for the integer divisibility and `native Decide` infrastructure. We are grateful to the constructive reverse mathematics community—especially the foundational work of Bishop, Bridges, Richman, and Ishihara—for developing the framework that makes calibrations like these possible.

The Lean 4 formalization was produced using AI code generation (Claude Code, Opus 4.6) under human direction. The author is a practicing cardiologist rather than a professional logician or arithmetic geometer; all mathematical claims should be evaluated on their formal content. We welcome constructive feedback from domain experts.

References

- [1] L. Berger, *An introduction to the theory of p -adic representations*, in: Geometric Aspects of Dwork Theory, de Gruyter, 2004, 255–292.
- [2] E. Bishop and D. Bridges, *Constructive Analysis*, Springer, 1985.

- [3] P. Colmez and J.-M. Fontaine, *Construction des représentations p -adiques semi-stables*, Invent. Math. **140** (2000), 1–43.
- [4] G. Faltings, *p -adic Hodge theory*, J. Amer. Math. Soc. **1** (1988), 255–299.
- [5] B. Mazur, J. Tate, and J. Teitelbaum, *On p -adic analogues of the conjectures of Birch and Swinnerton-Dyer*, Invent. Math. **84** (1986), 1–48.
- [6] J. H. Silverman, *The Arithmetic of Elliptic Curves*, 2nd ed., Graduate Texts in Mathematics **106**, Springer, 2009.
- [7] T. Tsuji, *p -adic étale cohomology and crystalline cohomology in the semi-stable reduction case*, Invent. Math. **137** (1999), 233–411.
- [8] P. C.-K. Lee, *Paper 45: Constructive Reverse Mathematics and Physics — Series Overview*, Zenodo, 2026. <https://doi.org/10.5281/zenodo.18676170>
- [9] P. C.-K. Lee, *Paper 50: Decidability landscape for the Standard Conjectures on abelian varieties*, Zenodo, 2026. <https://doi.org/10.5281/zenodo.18705837>
- [10] P. C.-K. Lee, *Paper 51: The Archimedean rescue — BSD calibration*, Zenodo, 2026. <https://doi.org/10.5281/zenodo.18732168>
- [11] P. C.-K. Lee, *Paper 54: The Bloch–Kato calibration*, Zenodo, 2026. <https://doi.org/10.5281/zenodo.18732964>
- [12] P. C.-K. Lee, *Paper 55: K3 surfaces, the Kuga–Satake construction, and the DPT framework*, Zenodo, 2026. <https://doi.org/10.5281/zenodo.18733731>
- [13] P. C.-K. Lee, *Paper 56: Self-Intersection of Exotic Weil Classes and Field Discriminants*, Zenodo, 2026. <https://doi.org/10.5281/zenodo.18734021>
- [14] P. C.-K. Lee, *Paper 57: Exotic Weil Self-Intersection Across All Nine Heegner Fields*, Zenodo, 2026. <https://doi.org/10.5281/zenodo.18735172>
- [15] P. C.-K. Lee, *Paper 58: Class Number Correction for Exotic Weil Classes on CM Abelian Fourfolds*, Zenodo, 2026. <https://doi.org/10.5281/zenodo.18734718>
- [16] S. Bloch and K. Kato, *L-functions and Tamagawa numbers of motives*, in: The Grothendieck Festschrift, vol. I, Birkhäuser, 1990, 333–400.
- [17] B. H. Gross and D. B. Zagier, *Heegner points and derivatives of L-series*, Invent. Math. **84** (1986), 225–320.
- [18] J. H. Silverman, *The difference between the Weil height and the canonical height on elliptic curves*, Math. Comp. **55** (1990), 723–743.
- [19] D. G. Northcott, *An inequality in the theory of arithmetic on algebraic varieties*, Proc. Cambridge Philos. Soc. **45** (1949), 502–509.
- [20] A. Borel, *Stable real cohomology of arithmetic groups*, Ann. Sci. Éc. Norm. Sup. **7** (1974), 235–272.
- [21] V. A. Kolyvagin, *Finiteness of $E(\mathbb{Q})$ and $\text{III}(E, \mathbb{Q})$ for a subclass of Weil curves*, Izv. Akad. Nauk SSSR **52** (1988), 522–540.